

# Fast Extended GCD for Large Integers for Verifiable Delay Functions

Kavya Sreedhar  
AHA Meeting 8.5.2021

# Advanced cryptography needs faster GCD


- GCD is a well-established operation in number theory with
  - Wide-ranging applications in cryptography
  - A focus on developing fast algorithms in the 1980s/90s
- Verifiable delay functions (VDFs), introduced in 2018, create a new need for understanding how fast work can be in this space

**Our goal: build fast GCD for the VDF application domain**


Verifiable delay functions

# Verifiable delay function (VDF)

- Function that is slow to compute but easy to verify
- Defined by 3 algorithms: *setup*, *evaluate*, *verify*
  - *Evaluate* takes a fixed number of sequential steps  $T$
  - *Verify* must be efficient when checking the result of *evaluate*



**Delay:** Requires minimum evaluation time despite any available parallelism



**Verifiable:** Can be efficiently verified

# The crypto community is excited about VDFs

## Chia Network Announces 2nd VDF Competition with \$100,000 in Total Prize Money

Matt Howard and Bram Cohen — April 4, 2019



THOMAS SIMMS

APR 22, 2019

## Protocol Labs and Ethereum Foundation Team Up to Research Verifiable Delay Functions



Protocol Labs

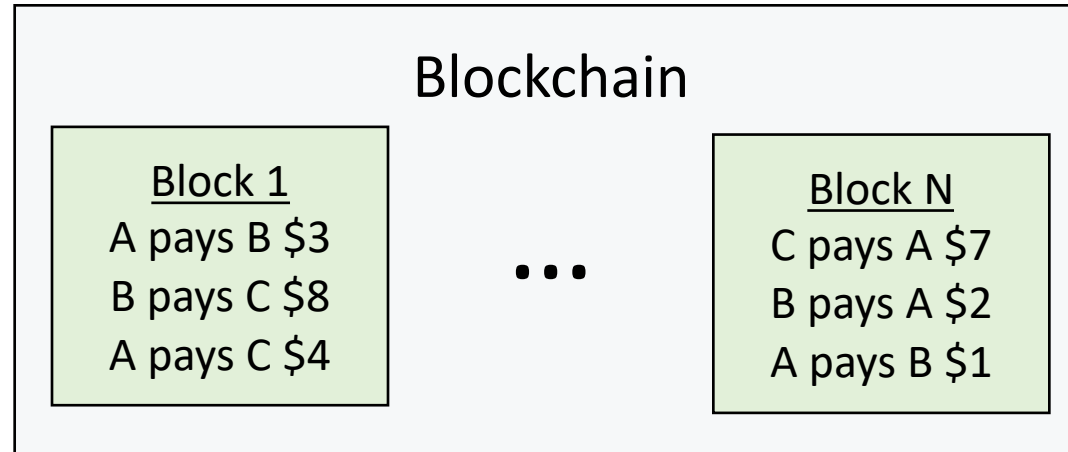


ethereum  
foundation

Updated Aug 11, 2019 at 6:57 p.m. PDT

At the cutting edge of blockchain research is a potential \$15 million dollar venture by the Ethereum Foundation centered around a technology called Verifiable Delay Functions (VDFs).

# VDFs can secure blockchain systems




- Transactions are recognized once validated and added to the chain
  - Verifiers validate transactions (and usually receive a reward for doing so)
  - Secure blockchains prevent attackers from taking over a majority of validation

# VDFs can secure blockchain systems

*time* →

Challenge  
chain




1. Challenges broadcast to the network

---

Infused  
challenge  
chain

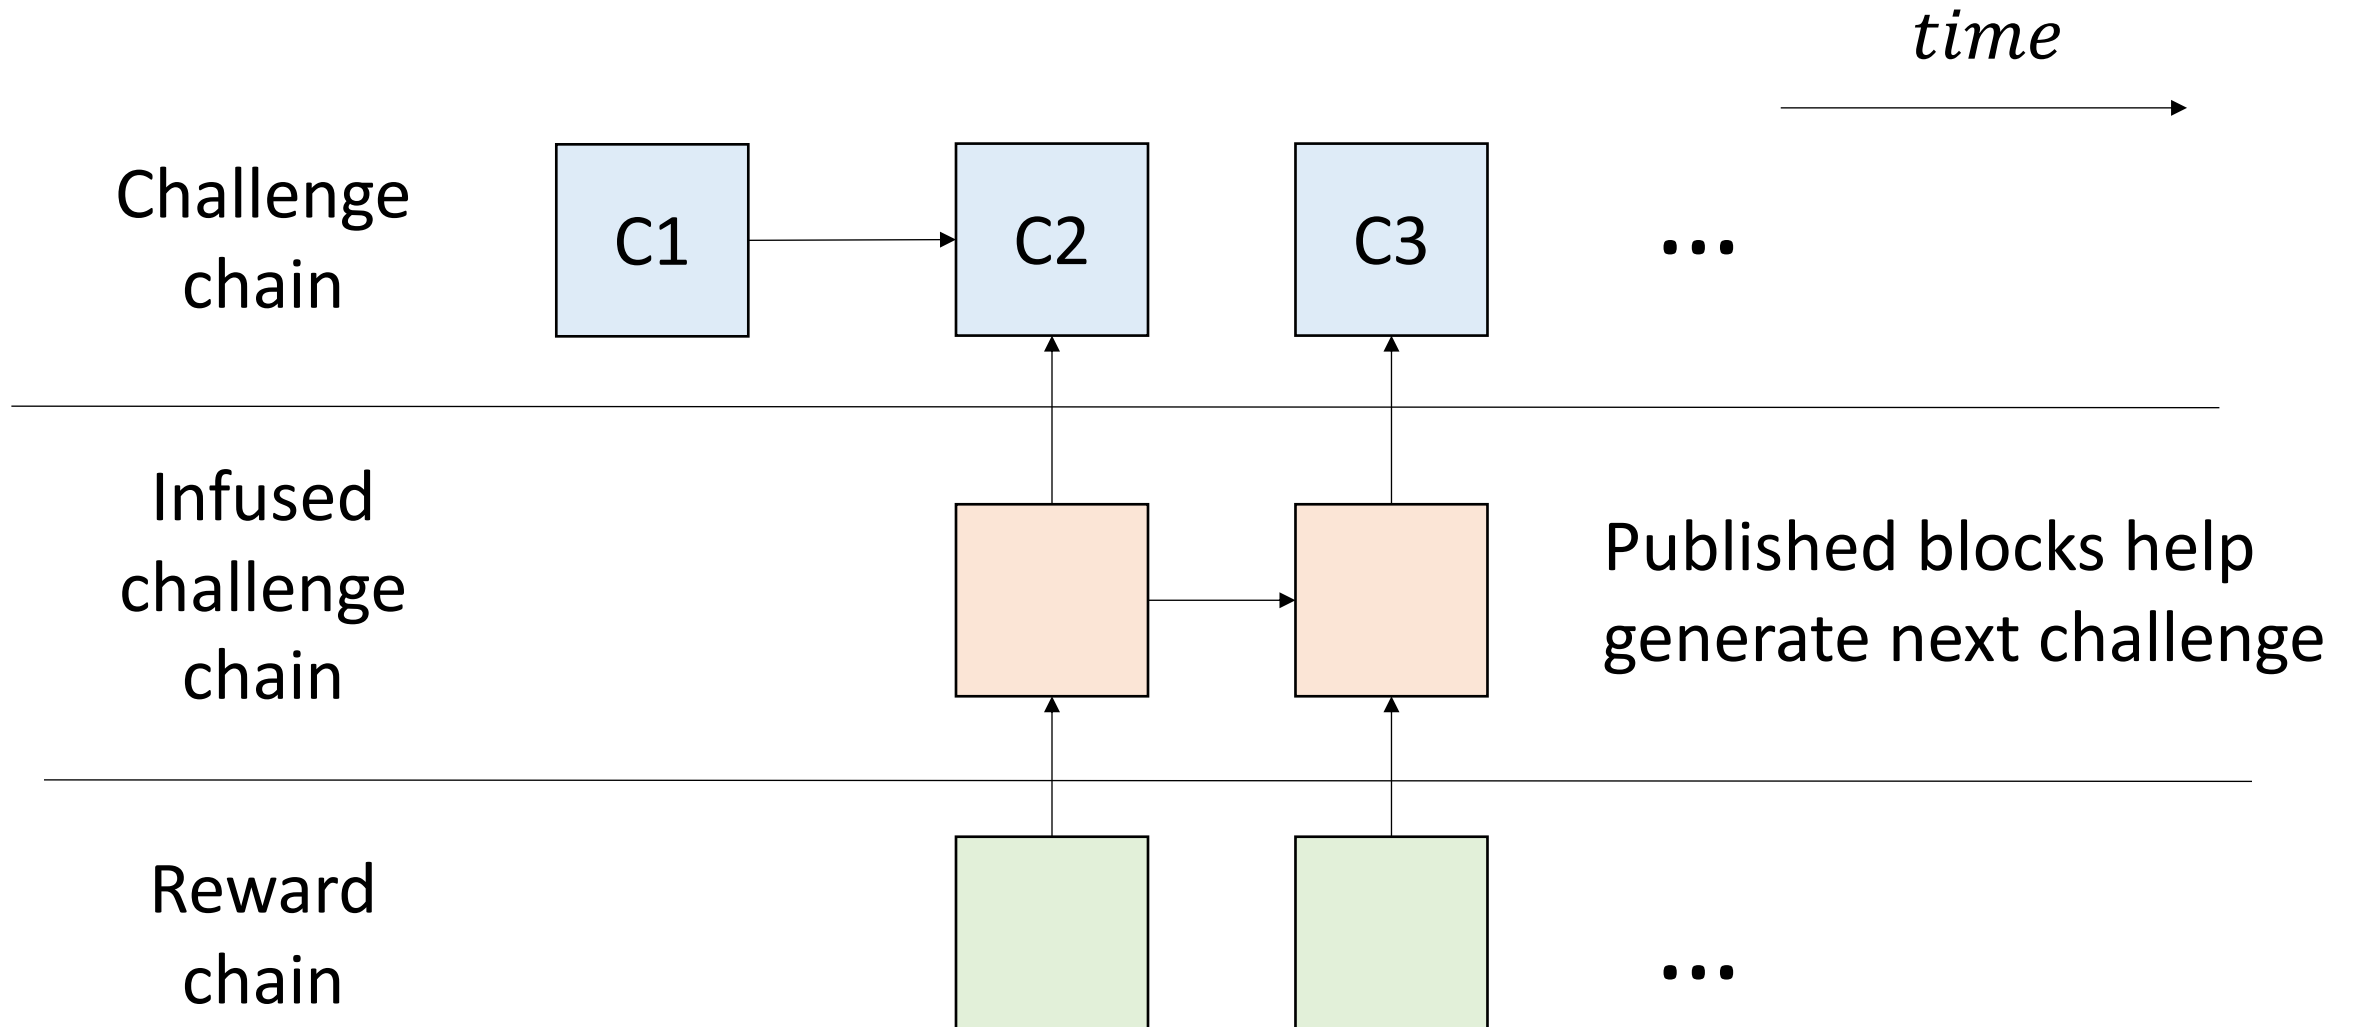
---

Reward  
chain



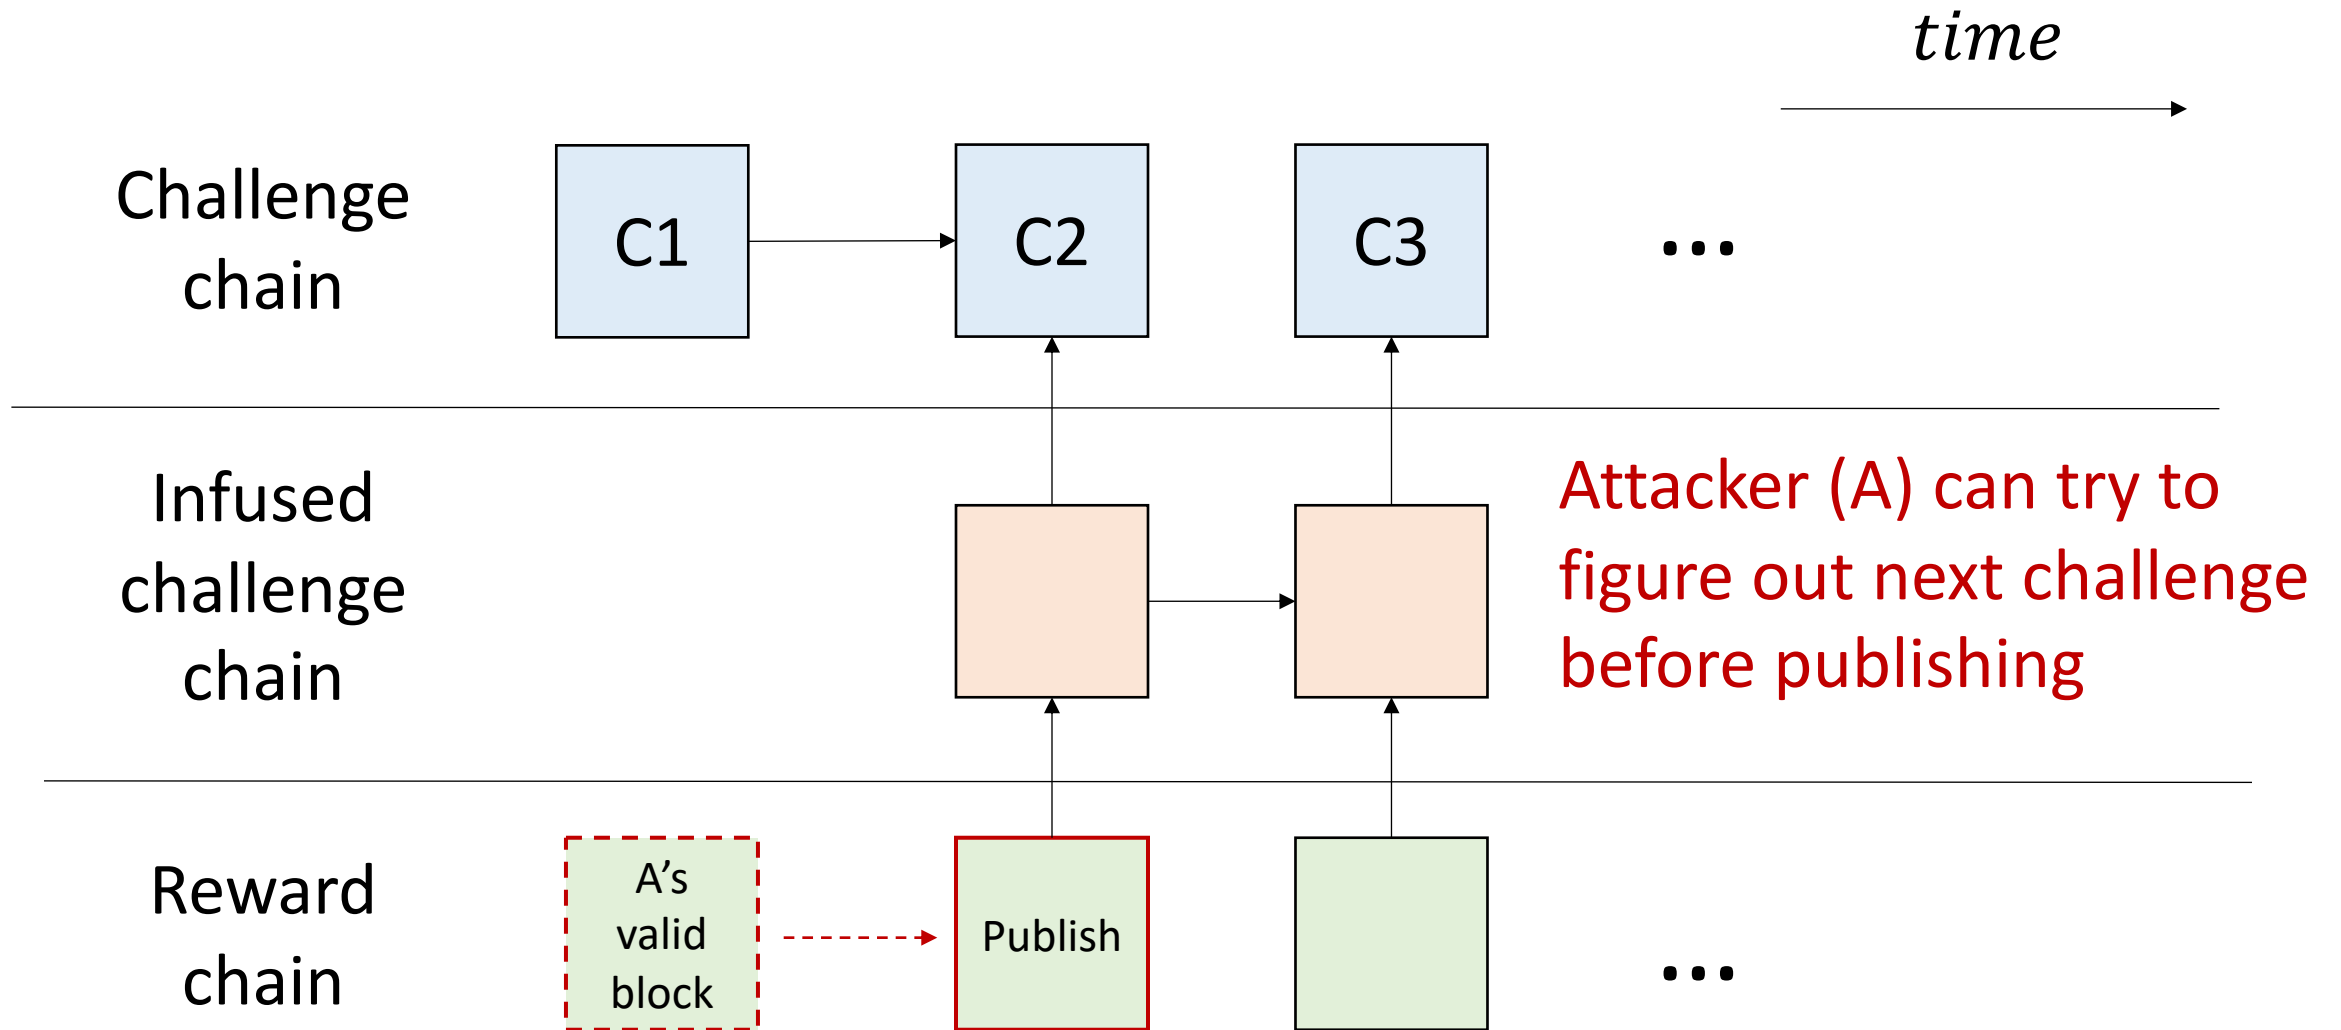
2. Parties see if they have a valid  
block and publish if so

# VDFs can secure blockchain systems

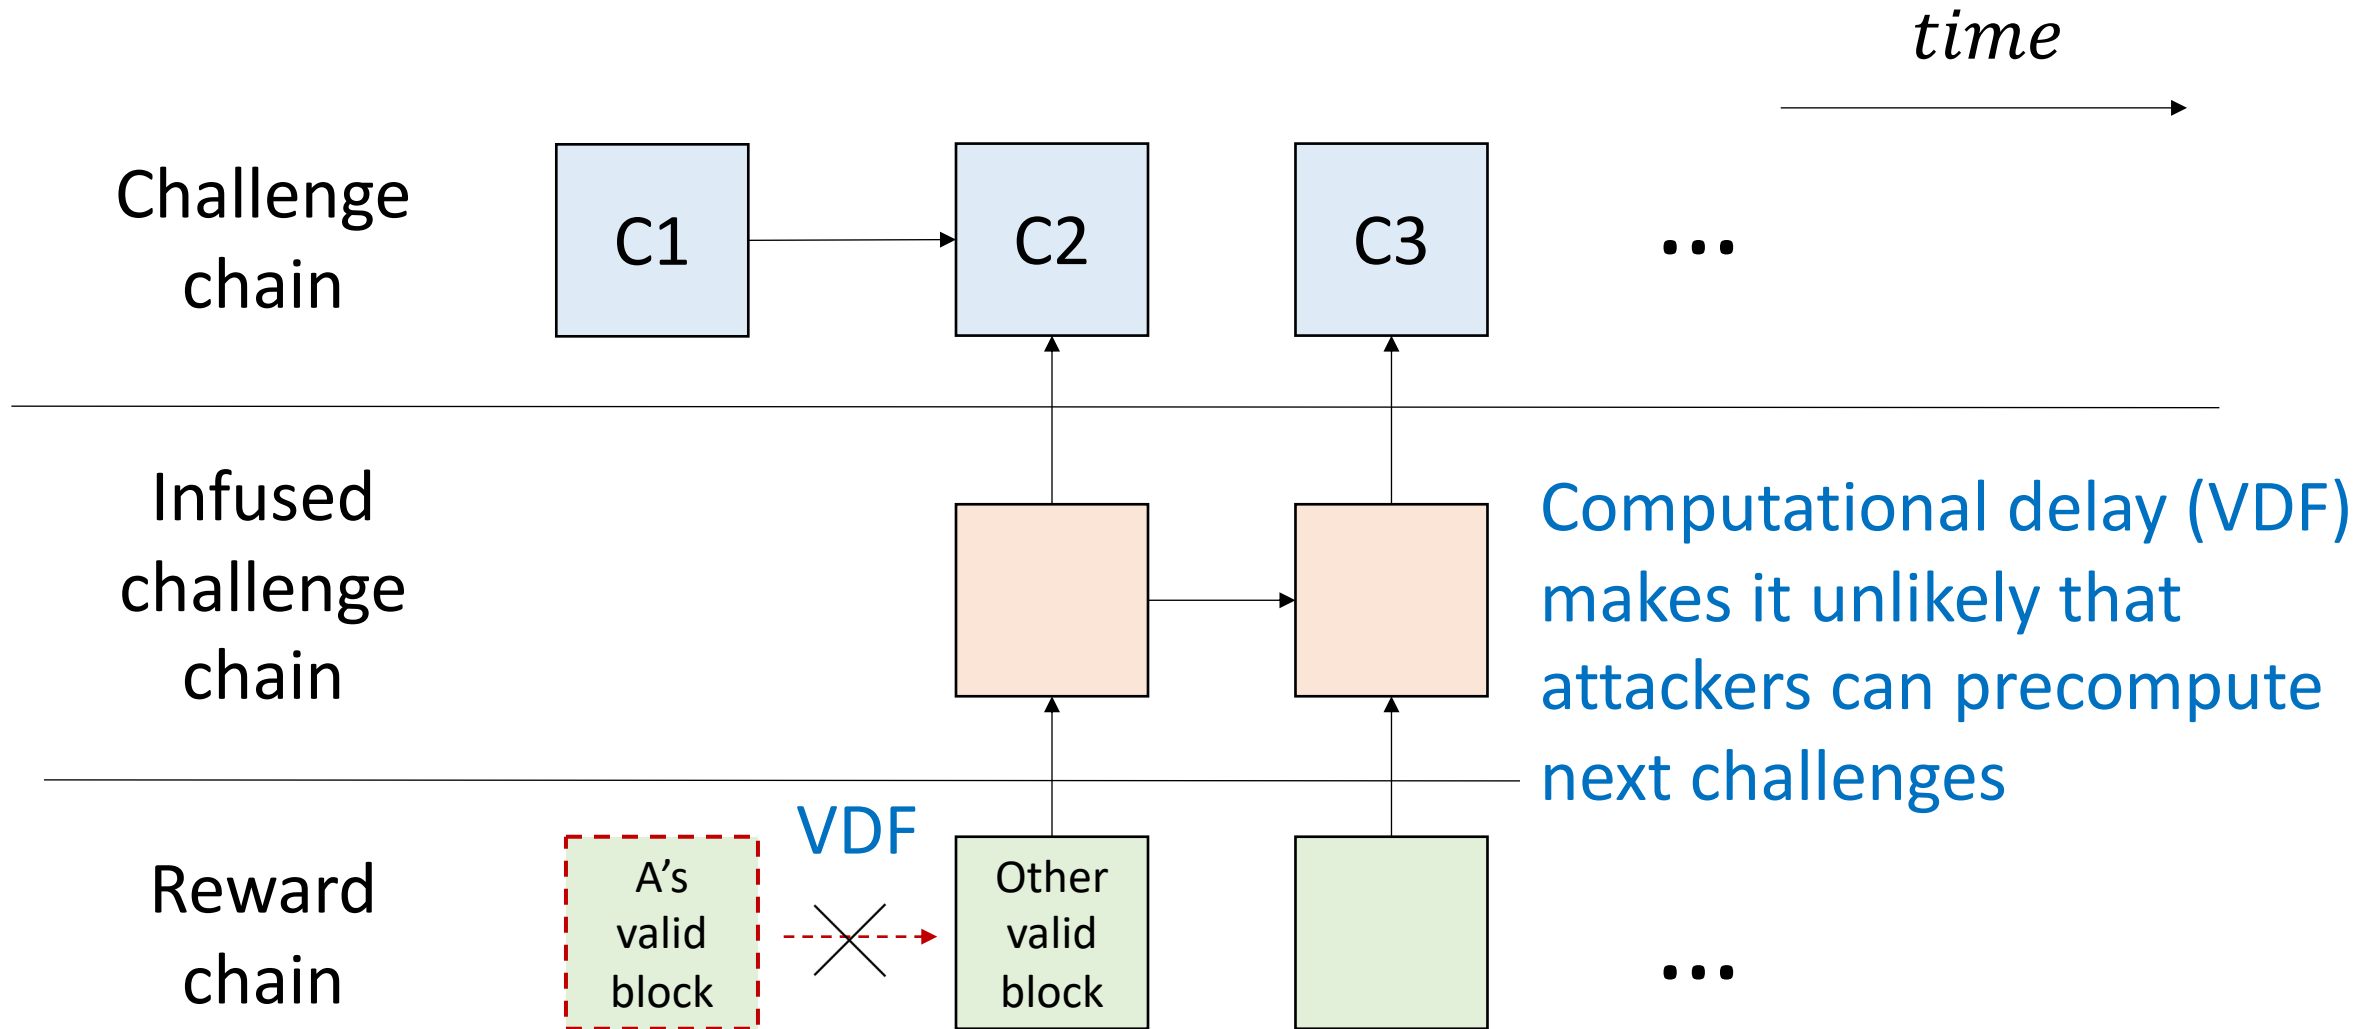




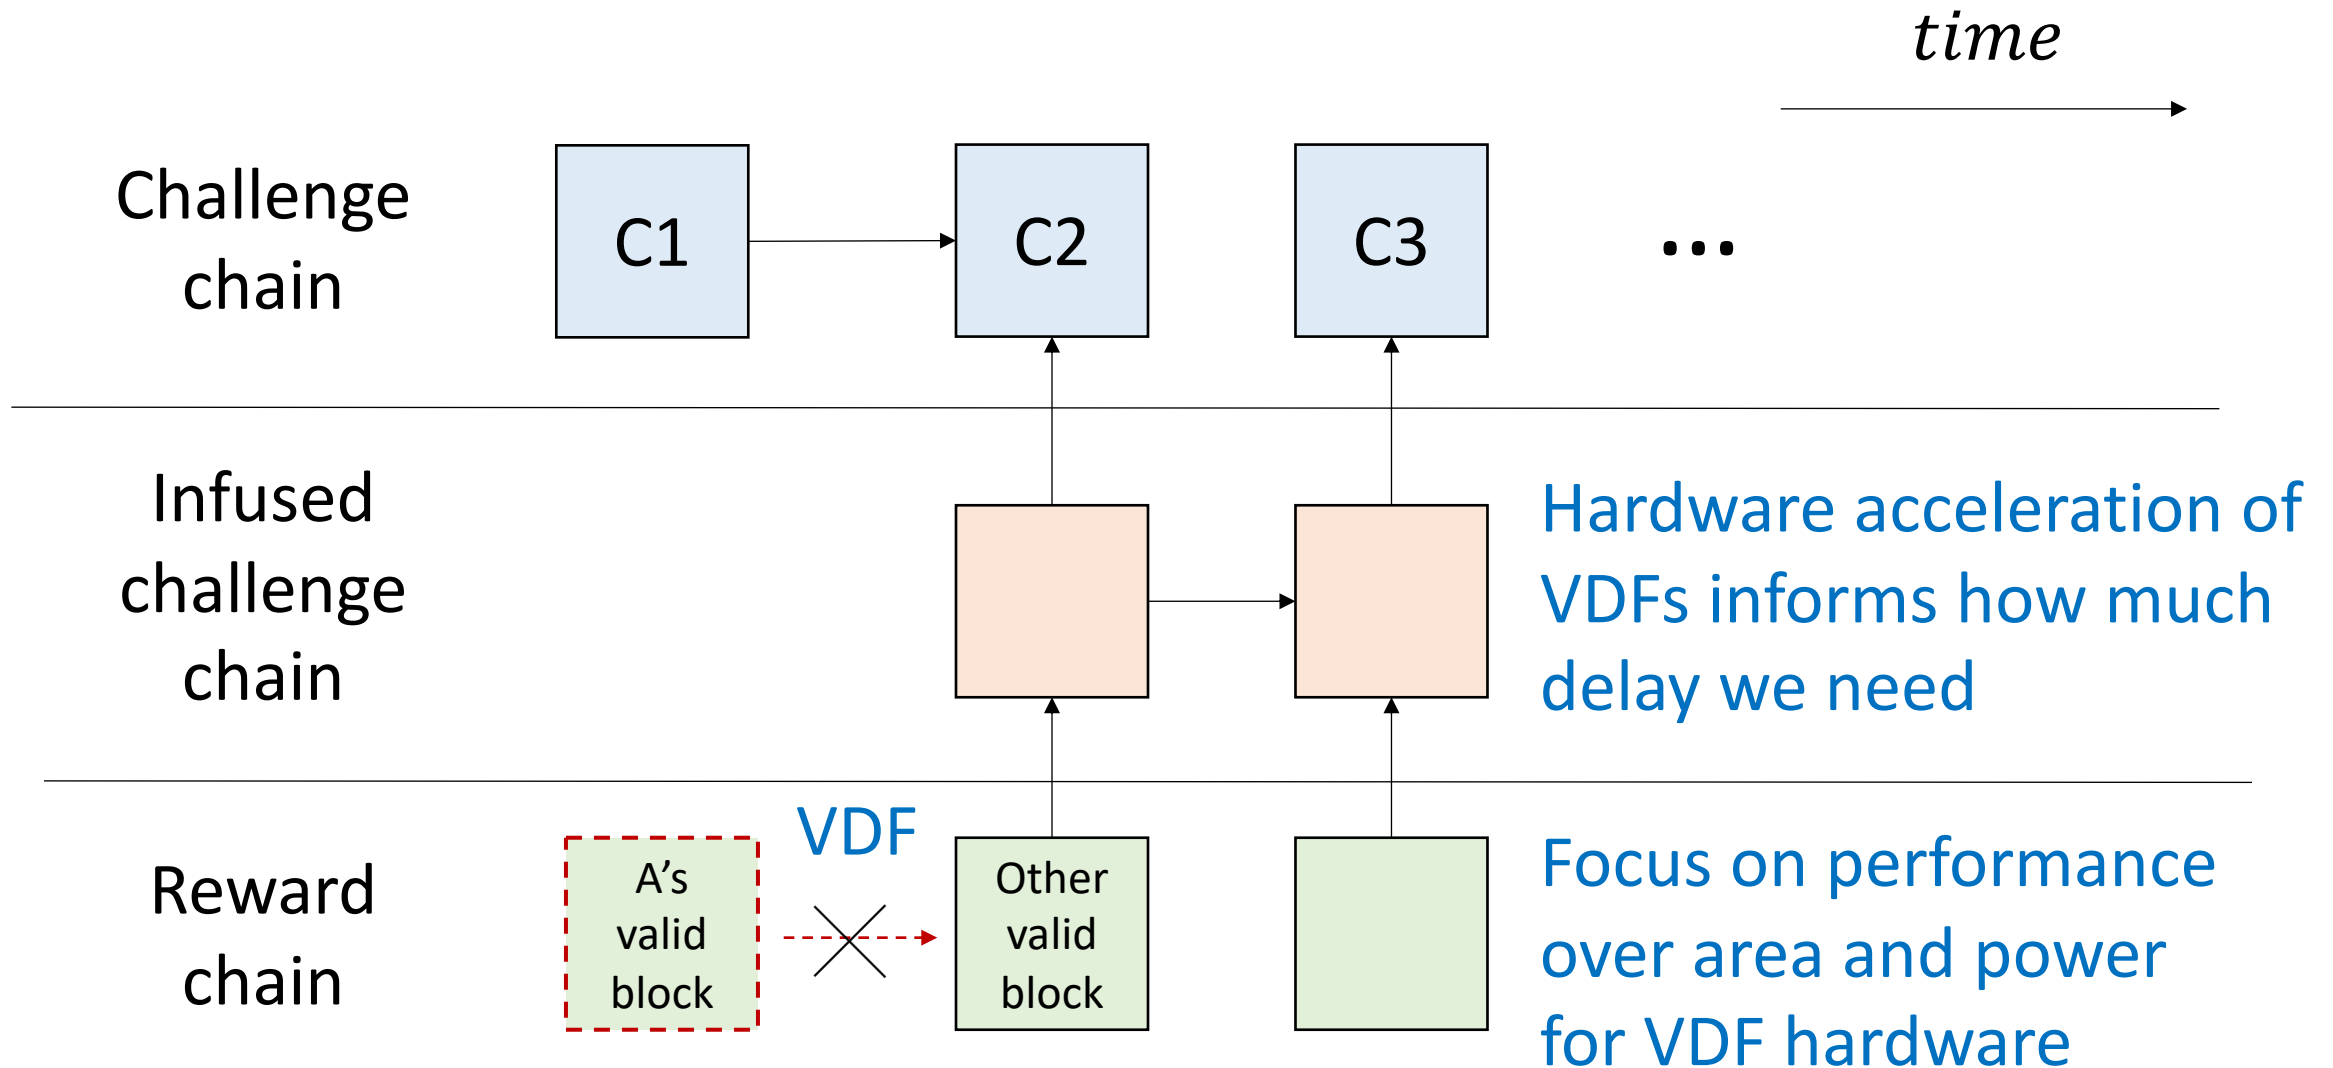
# VDFs can secure blockchain systems



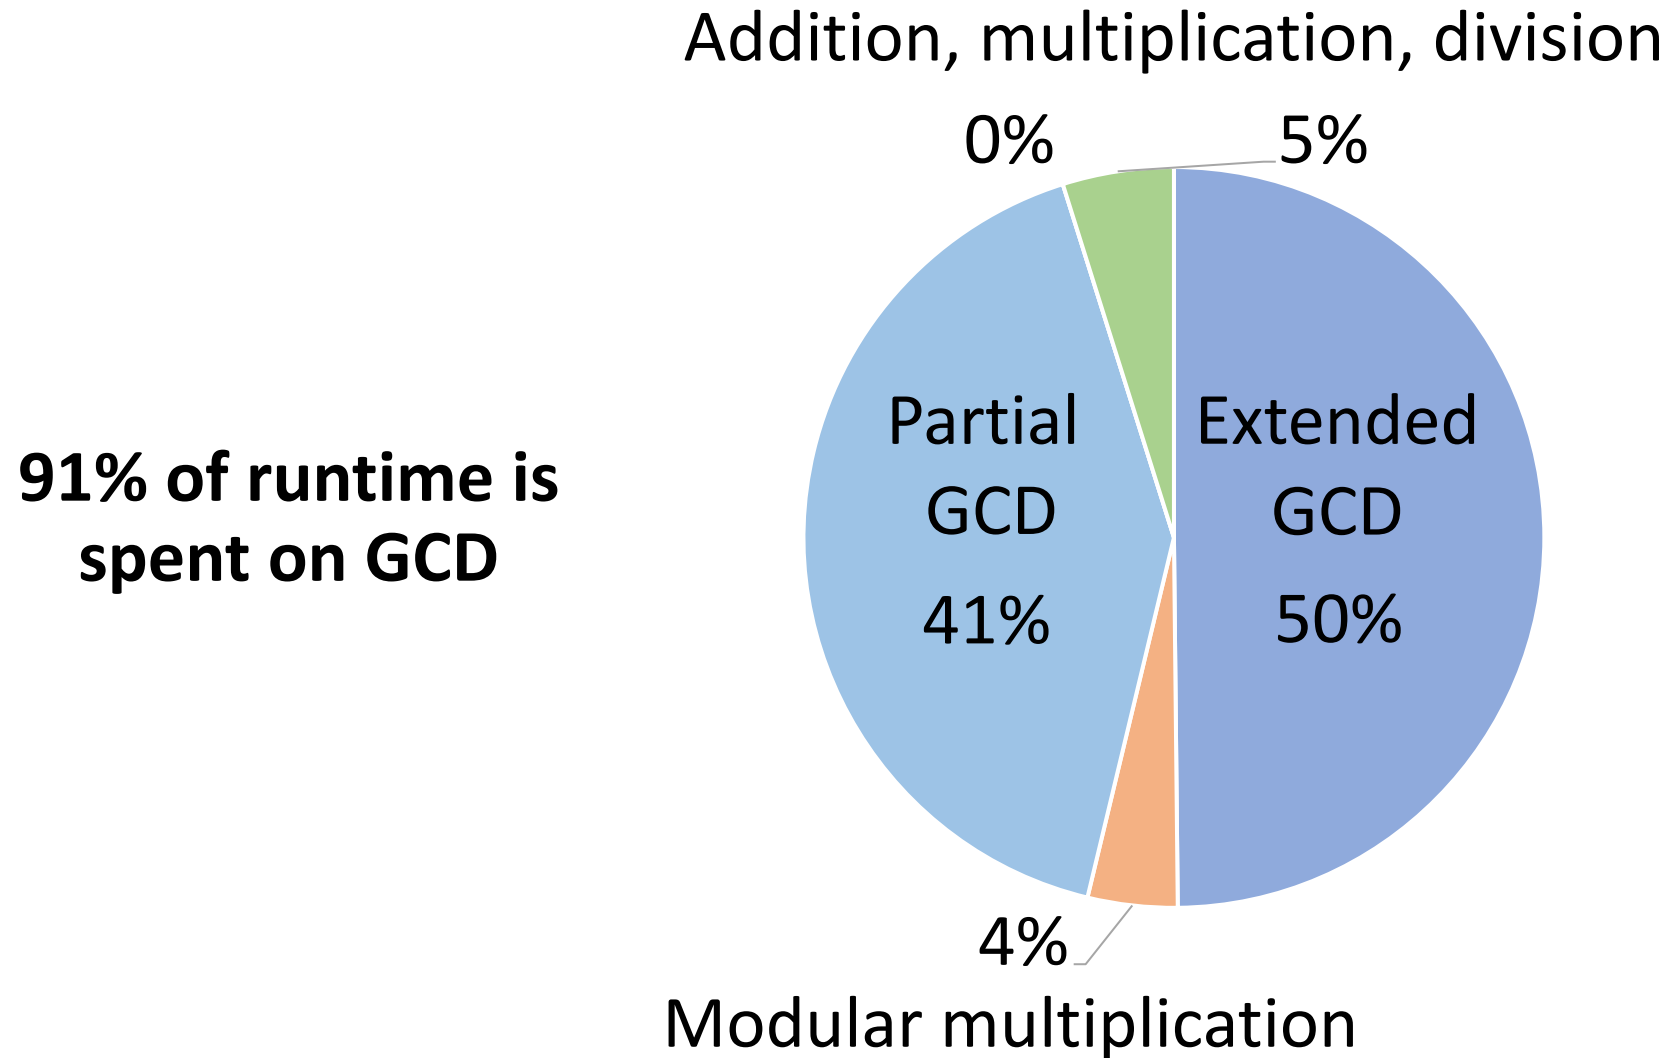
# VDFs can secure blockchain systems



# VDFs can secure blockchain systems



# VDFs primarily rely on fast extended GCDs



# GCD Algorithms

Algorithms use GCD-preserving transformations

$$\mathbf{g} = \mathbf{gcd}(\mathbf{a}, \mathbf{b}) = \mathbf{gcd}(\mathbf{a} - \mathbf{b}, \mathbf{b})$$

$$\mathbf{a} = \mathbf{g} * \mathbf{a_g}, \quad \mathbf{b} = \mathbf{g} * \mathbf{b_g}$$

$$\Rightarrow \mathbf{a} - \mathbf{b} = \mathbf{g} * (\mathbf{a_g} - \mathbf{b_g})$$

$$3 = \mathbf{gcd}(33, 9) = \mathbf{gcd}(24, 9)$$

Algorithms use GCD-preserving transformations

$$\mathbf{g = gcd(a, b) = gcd(a - b, b)}$$

$$a = g * a_g, \quad b = g * b_g$$

$$\Rightarrow a - b = g * (a_g - b_g)$$

$$3 = gcd(33, 9) = gcd(24, 9)$$

$$\mathbf{gcd(a, b) = gcd(a \bmod b, b)}$$

$$a = g * a_g, \quad b = g * b_g$$

$$\Rightarrow a \bmod b = a - b * q = g * (a_g - b_g * q)$$

$$3 = gcd(33, 9) = gcd(6, 9)$$

# Algorithms use GCD-preserving transformations

**Steins**     $g = \gcd(a, b) = \gcd(a - b, b)$

$$a = g * a_g, \quad b = g * b_g$$

$$\Rightarrow a - b = g * (a_g - b_g)$$

$$3 = \gcd(33, 9) = \gcd(24, 9)$$

**Euclid**     $\gcd(a, b) = \gcd(a \bmod b, b)$

$$a = g * a_g, \quad b = g * b_g$$

$$\Rightarrow a \bmod b = a - b * q = g * (a_g - b_g * q)$$

$$3 = \gcd(33, 9) = \gcd(6, 9)$$



# GCD algorithms example: $G(27, 2) = 1$

Euclid:  $\gcd(a \bmod b, b)$

a	b	Operation
27	2	original a, b
2	1	$27 \bmod 2 = 1$
<b>1</b>	0	$2 \bmod 1 = 0$

Euclid's converges faster, but Steins uses simpler operations (shifts, adds, comparisons) instead of division

Steins: shift if even or else  $\gcd(a - b, b)$

a	b	Operation
27	2	original a, b
27	1	$b / 2$
26	1	subtract
13	1	$a / 2$
12	1	subtract
6	1	$a / 2$
3	1	$a / 2$
2	1	subtract
1	1	$a / 2$
<b>1</b>	0	subtract

# Extended GCD

- Computes Bezout coefficients  $b_a, b_b$ :  $b_a * a_0 + b_b * b_0 = \gcd(a_0, b_0)$
- GCD algorithms extended to calculate  $b_a, b_b$  by maintaining these relations every cycle where  $\gcd(a_0, b_0) = \gcd(a, b)$ :

$$\begin{aligned}u * a_0 + m * b_0 &= a \\ y * a_0 + n * b_0 &= b\end{aligned}$$

- At start,  $u = n = 1$  and  $m = y = 0$
- At end,  $\gcd(a_0, b_0) = a$ ,  $b_a = u$ ,  $b_b = m$

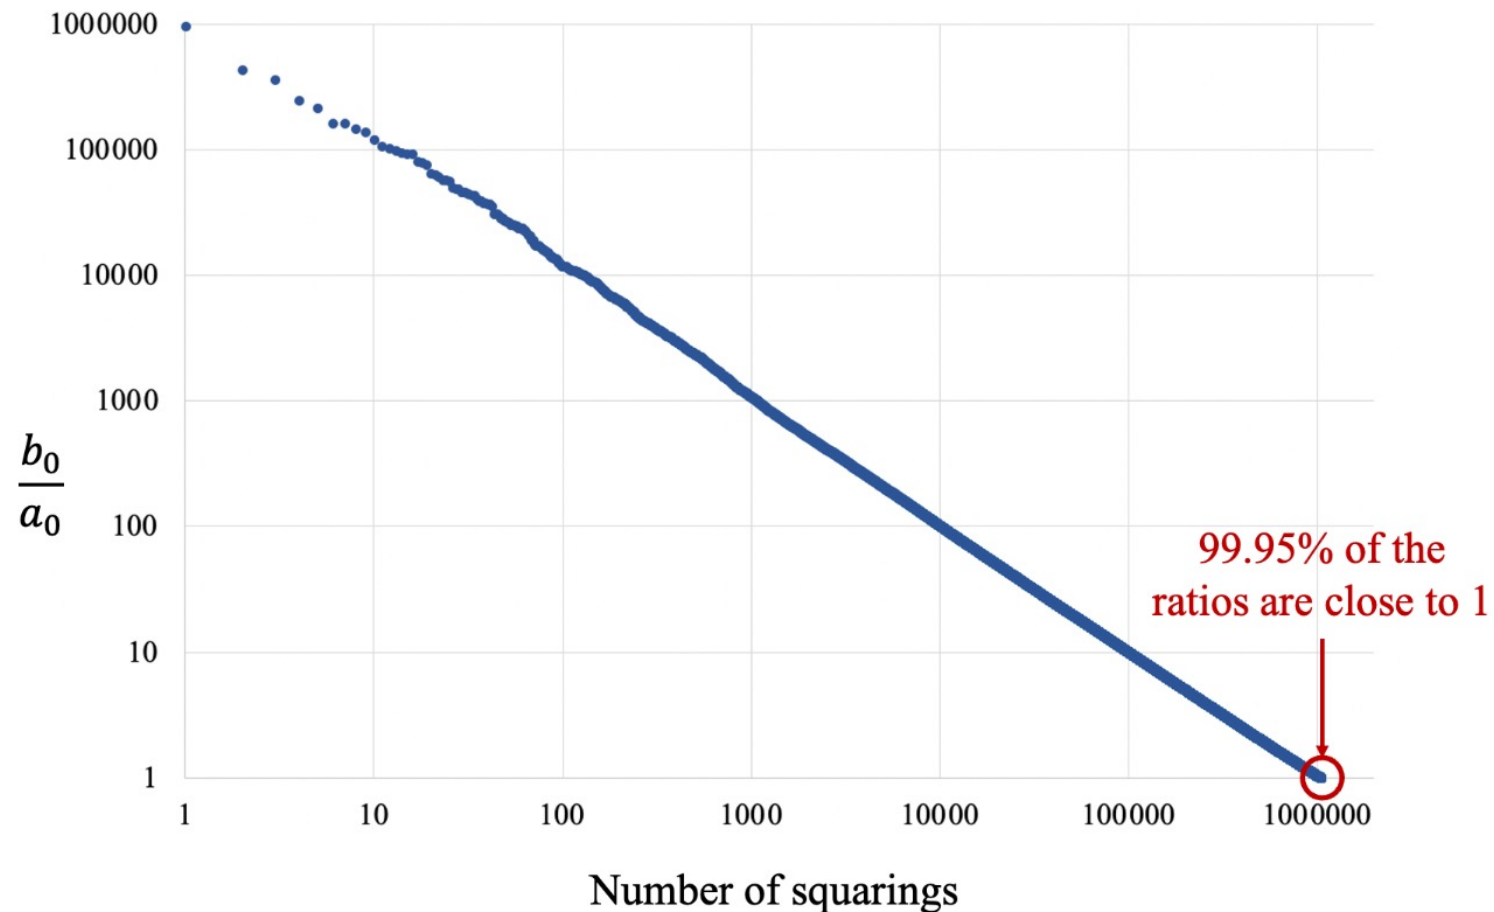
# GCD properties in the VDF context

- Bezout coefficients are needed
- Inputs are integers that are
  - Large – 1024+ bits
  - Relatively prime – GCD is known to be equal to 1
- Performance is primary hardware focus

**We aim to find state-of-the-art starting points from existing literature that are well-suited to these properties.**

# VDF inputs are well-suited to Stein's GCD

- Subtraction can reduce many bits when inputs are similar in length



# Hardware-friendly GCD algorithms

# Avoiding carry propagation for large-integer GCD

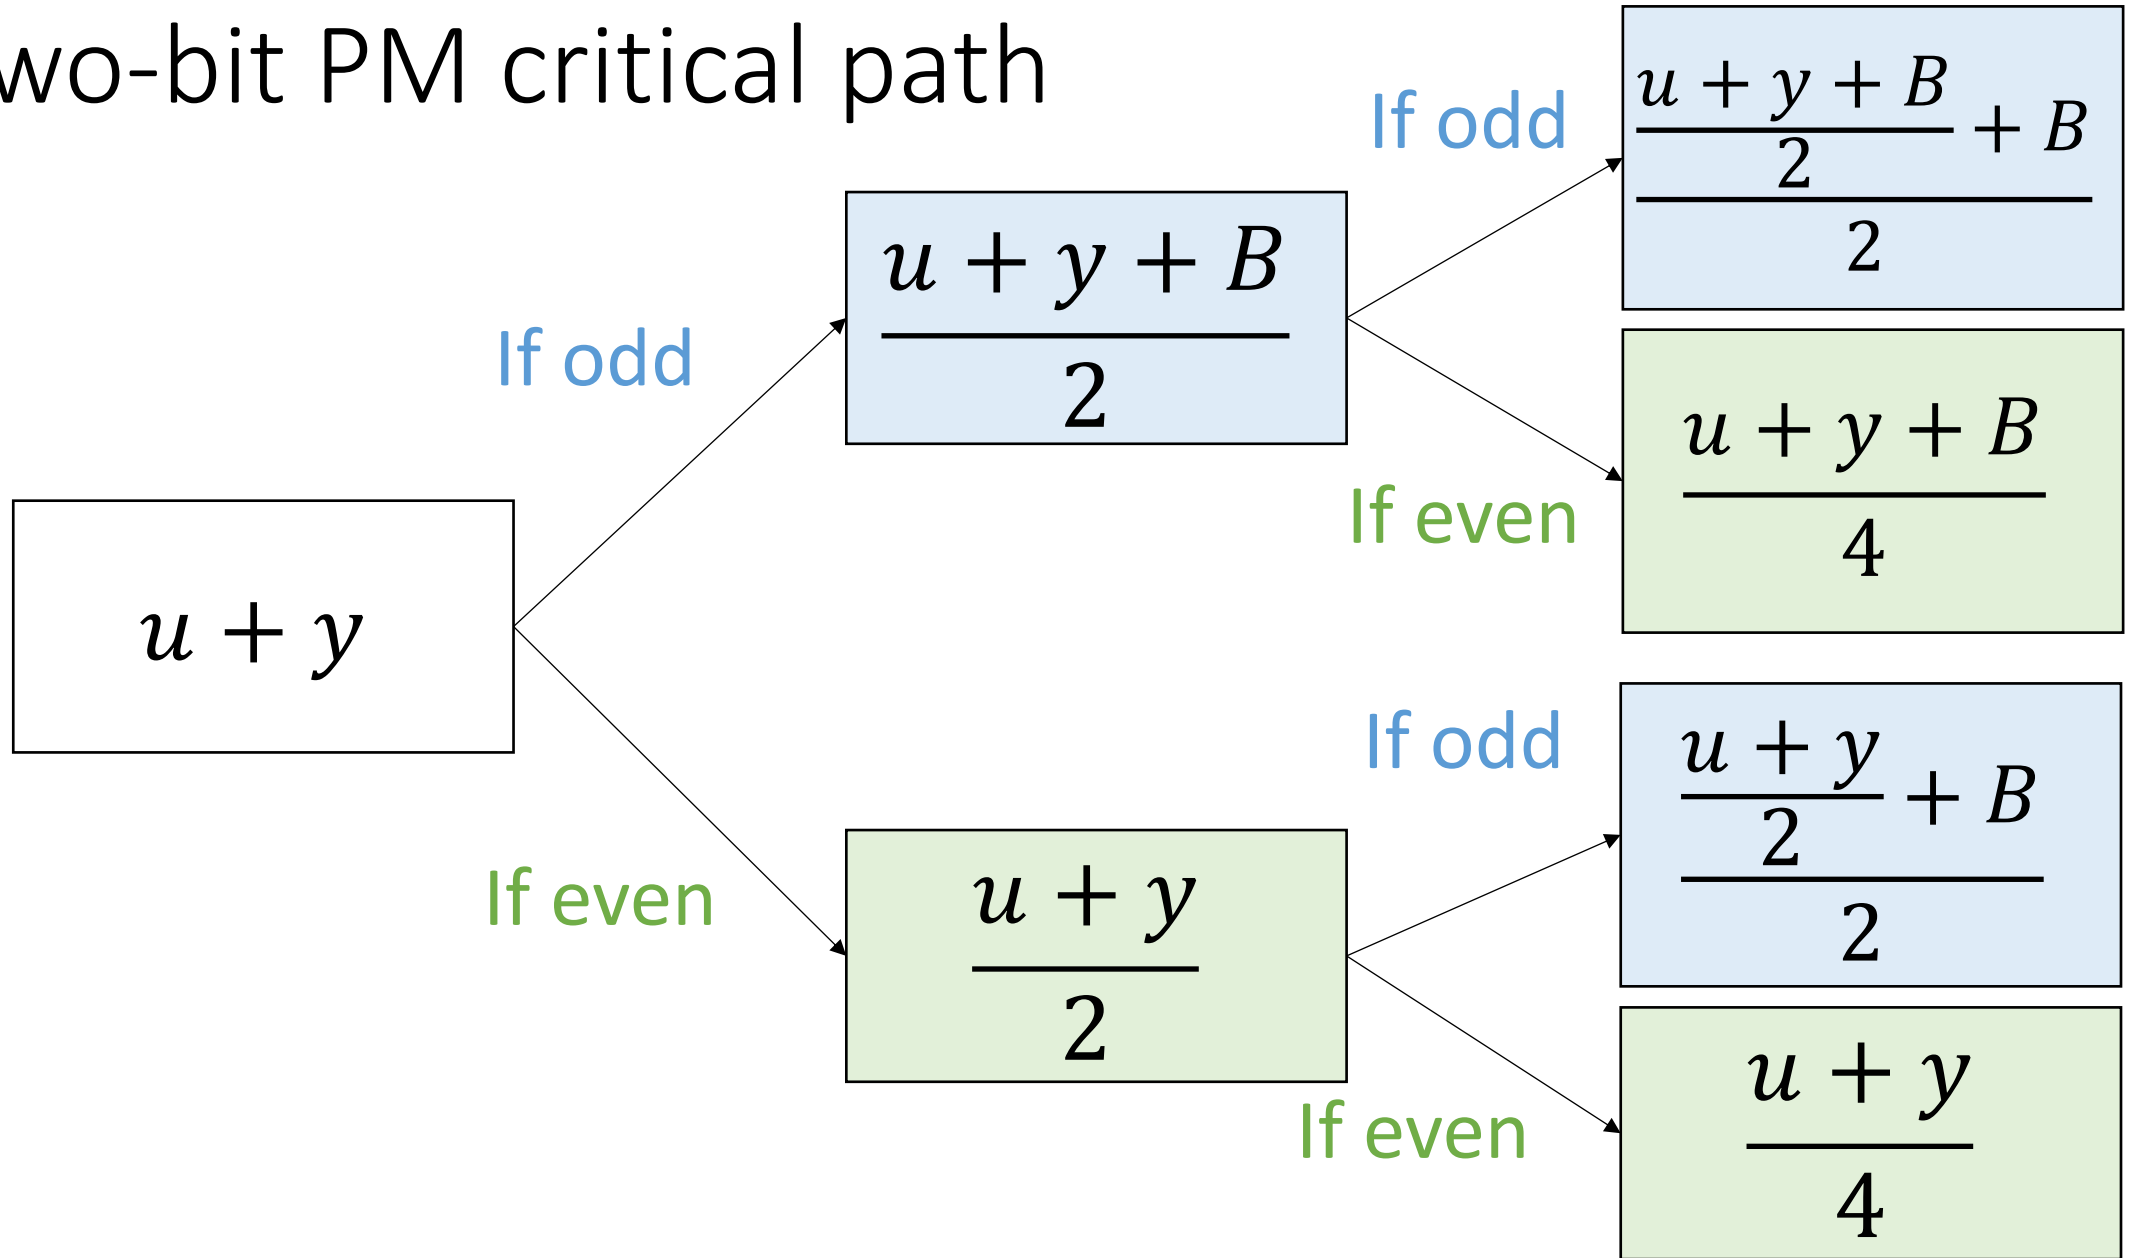
- Stein's algorithm has many back-to-back additions by design
- Carry-save adders (CSAs) eliminate carry propagation in these adds
  - $O(1)$  delay instead of  $O(\text{number of bits})$
- Algorithms avoid needing  $a, b$  to check  $a > b$  by
  - updating  $G(a, b) = G\left(\frac{a+b}{2}, \frac{a-b}{2}\right)$  when  $a, b$  odd
  - using approximate value of  $\log_2 a - \log_2 b$  to decide which branch to take
- Two-bit PM further reduces 2 bits if possible and avoids swapping  $a, b$

# Two-bit Plus-Minus (PM) algorithm

Case Number	Description
1	$a$ is divisible by 4
2	$a$ is divisible by 2 (and not by 4)
3	$b$ is divisible by 4
4	$b$ is divisible by 2 (and not by 4)
5	$\delta \geq 0$ and $a + b$ is divisible by 4
6	$\delta \geq 0$ and $b - a$ is divisible by 4
7	$\delta < 0$ and $a + b$ is divisible by 4
8	$\delta < 0$ and $b - a$ is divisible by 4


- Bits reduced per cycle
- Worst case: 1
  - Best case:  $2 + \min(\log_2 a, \log_2 b)$
  - Reduces more bits than Stein's 80% of the time

# Two-bit PM critical path





# Two-bit PM critical path




- $u, y$  are in CSA form, so  $u + y$  requires 2 CSAs
- $B$  is not in CSA form, so critical path has **4 CSAs**


# Our GCD Design

Main Loop


GCD




Bézout

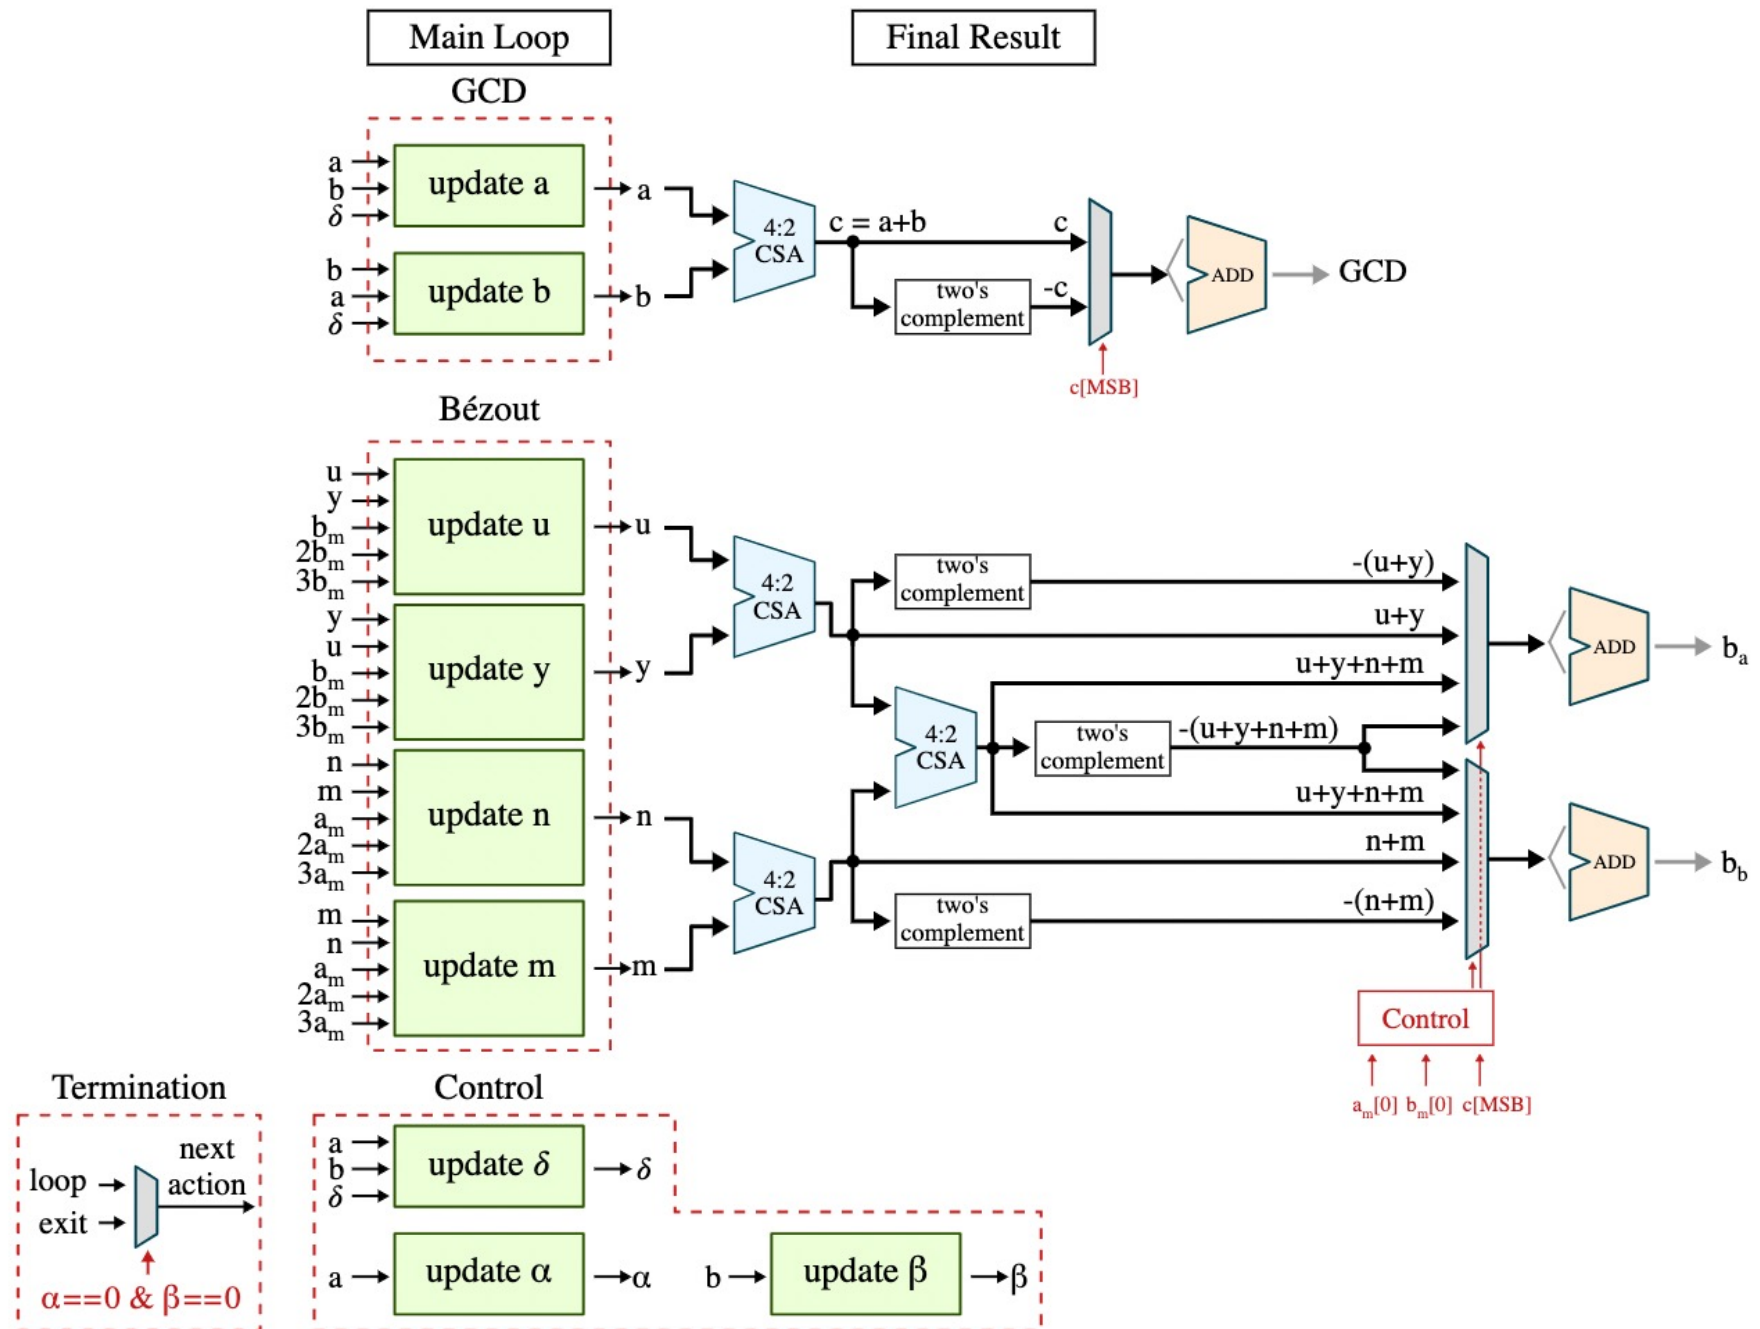


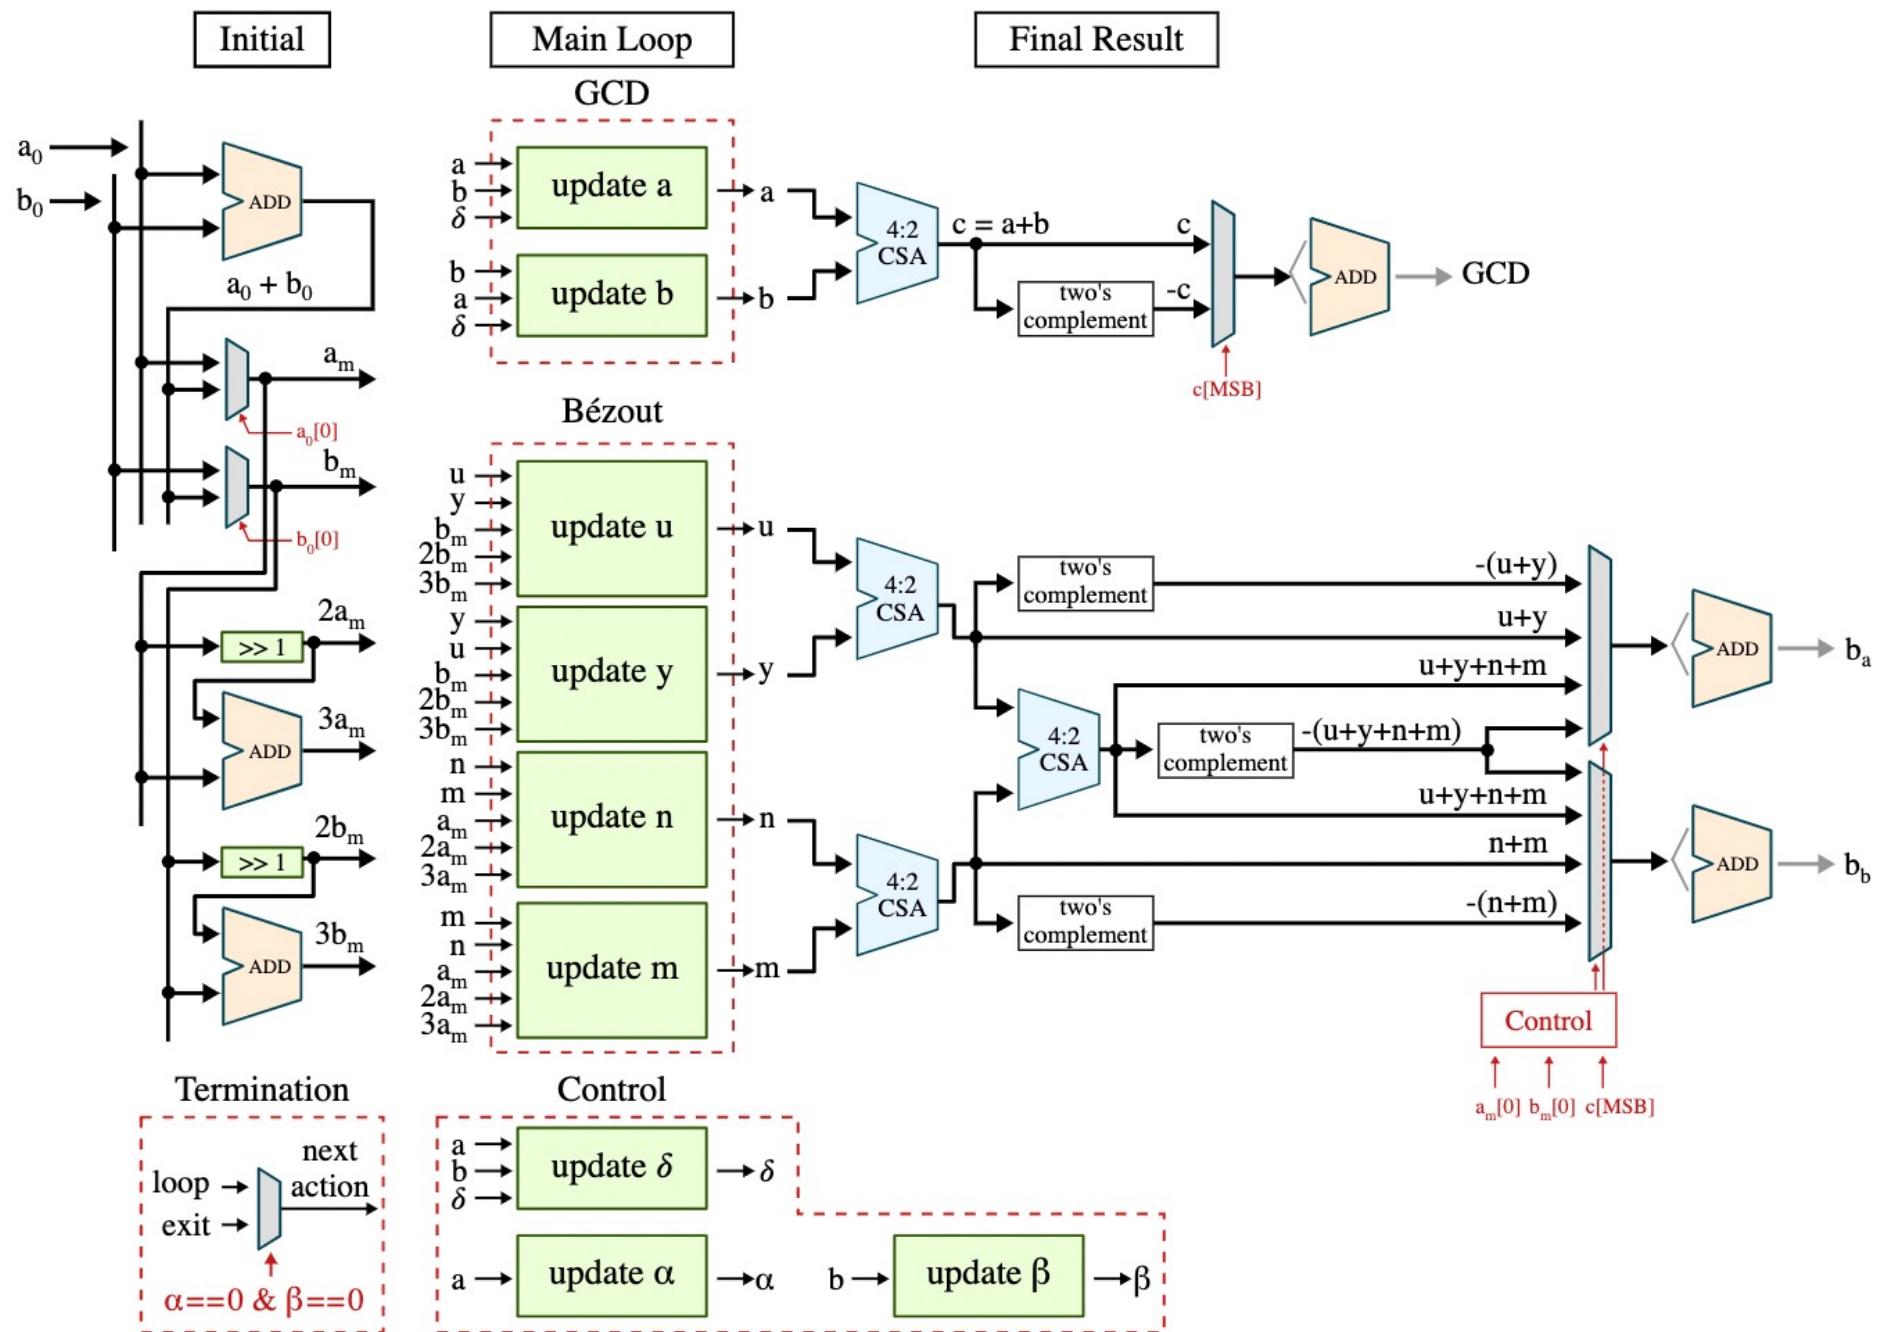
Termination

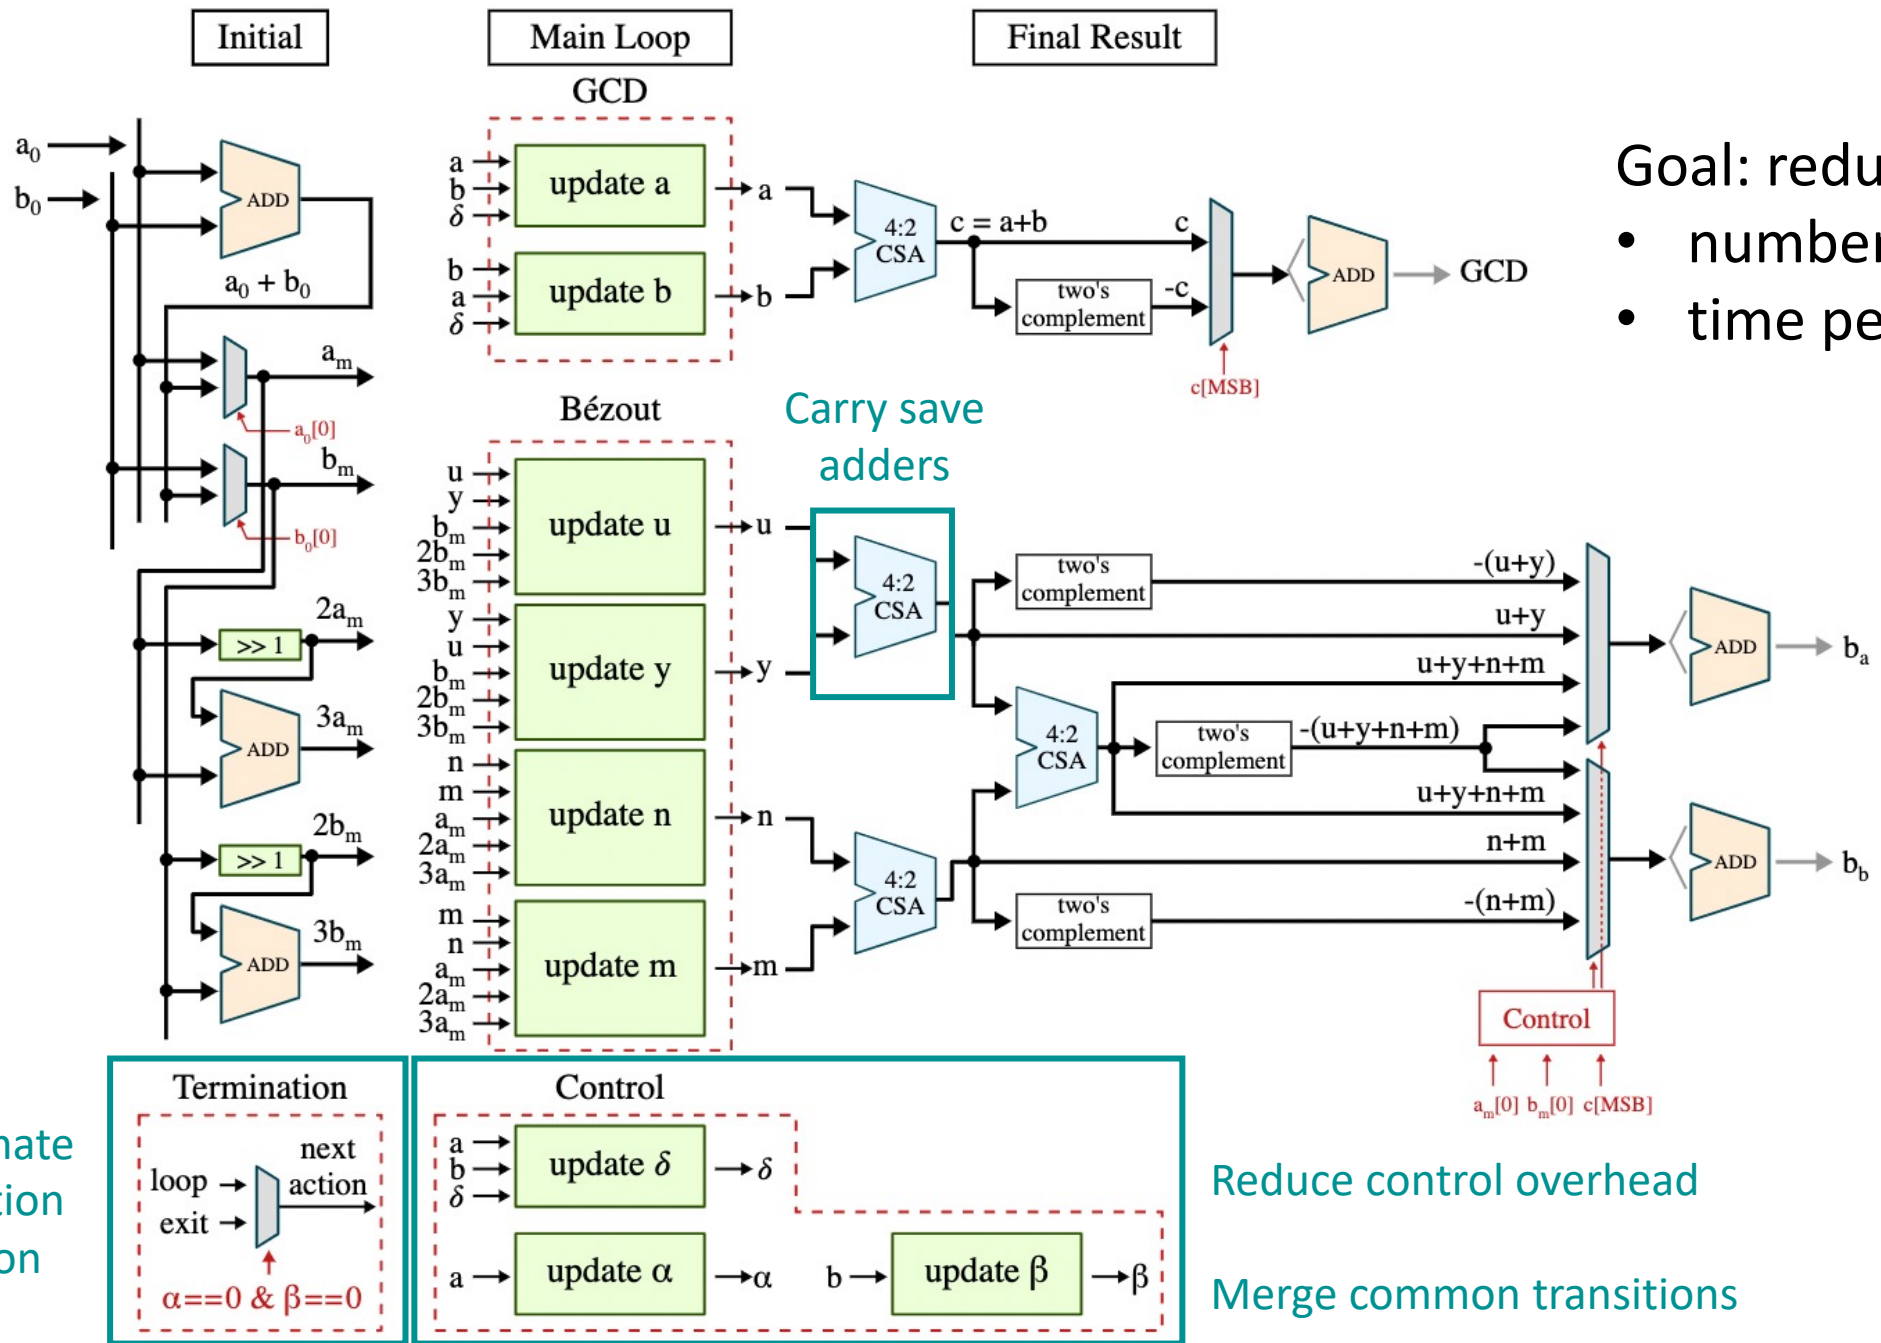


Control









Goal: reduce

- number of cycles
- time per cycle


Approximate  
termination  
condition

Reduce control overhead

Merge common transitions

# Terminate based on approximate binary logs


- We check if  $\alpha \cong \log_2 a, \beta \cong \log_2 b$  are 0 instead of  $a, b$  to
  - Avoid the carry-propagate add to get  $a, b$  from CSA form
- However, this is an approximation and adds 154 cycles
- So, we get  $a$  and  $b$  occasionally with carry-propagate adds to compute
$$\alpha = \log_2 a, \beta = \log_2 b$$
- This occasional correction reduces the overhead to 43 cycles (3.5%)



**Performance impact: Reduce 2 long 1024-bit carry-propagate checks to 2 faster 10-bit carry-propagate checks**



# Reducing cycle count by merging transitions




Two-bit PM Transition Matrix

	Description
1	a is divisible by 4
2	a is divisible by 2 (and not by 4)
3	b is divisible by 4
4	b is divisible by 2 (and not by 4)
5	$\delta \geq 0$ and $a + b$ is divisible by 4
6	$\delta \geq 0$ and $b - a$ is divisible by 4
7	$\delta < 0$ and $a + b$ is divisible by 4
8	$\delta < 0$ and $b - a$ is divisible by 4




# Merging common transitions



	Description
1	a is divisible by 4
2	a is divisible by 2 (and not by 4)
3	b is divisible by 4
4	b is divisible by 2 (and not by 4)
5	$\delta \geq 0$ and $a + b$ is divisible by 4
6	$\delta \geq 0$ and $b - a$ is divisible by 4
7	$\delta < 0$ and $a + b$ is divisible by 4
8	$\delta < 0$ and $b - a$ is divisible by 4

# Include divide by 8, 16 cases



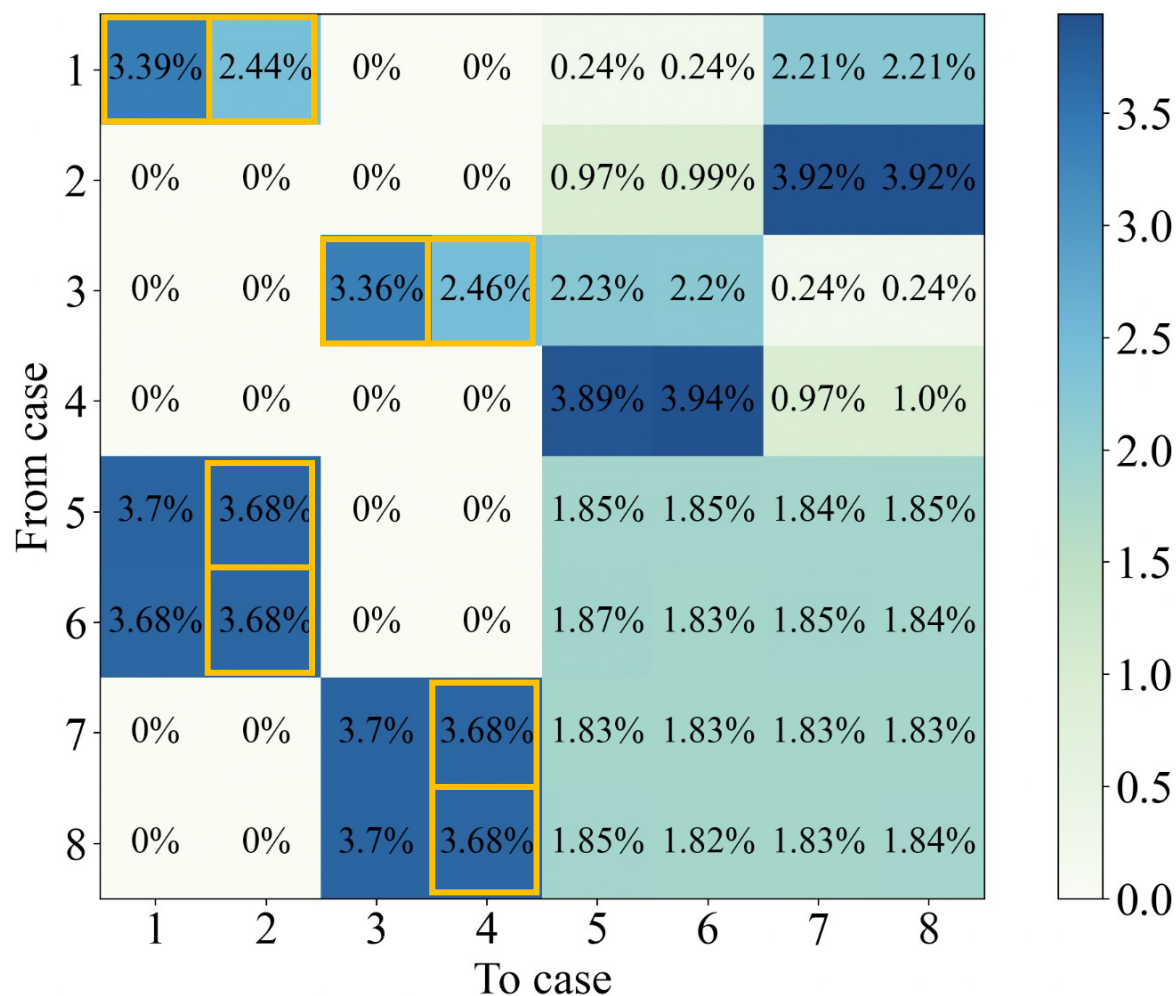
	Description
1	a is divisible by 4
2	a is divisible by 2 (and not by 4)
3	b is divisible by 4
4	b is divisible by 2 (and not by 4)
5	$\delta \geq 0$ and $a + b$ is divisible by 4
6	$\delta \geq 0$ and $b - a$ is divisible by 4
7	$\delta < 0$ and $a + b$ is divisible by 4
8	$\delta < 0$ and $b - a$ is divisible by 4

# Reducing cycle count summary

Transition Optimization	Cycles	Datapath (ns)	GCD Runtime (us)
Baseline	0	0	3.70
<b>Merge 5-2 into 5</b> <b>Merge 7-4 into 7</b>	<b>-236</b>	<b>+0.6 (1 CSA)</b>	<b>3.60</b>
Merge 5-2, 5-1 into 5 Merge 7-4, 7-3 into 7	-374	+1.2 (2 CSAs)	3.62
Merge 5-2, 5-1, 5-1-1, 5-1-2 Merge 7-4, 7-3, 7-3-3, 7-3-4	-464	+2.4 (4 CSAs)	4.17
Add cases: a, b divisible by 8	-65	0	3.52
<b>Add cases: a, b divisible by 16, 8</b>	<b>-92</b>	<b>0</b>	<b>3.44</b>
Bolded rows	-328	+0.6 (1 CSA)	3.27

**Performance impact: 11.6% faster runtime by trading a small increase in cycle time for lower cycle counts**

# Reducing cycle count chosen mergings



	Description
<b>1</b>	a is divisible by 4
<b>2</b>	a is divisible by 2 (and not by 4)
<b>3</b>	b is divisible by 4
<b>4</b>	b is divisible by 2 (and not by 4)
<b>5</b>	$\delta \geq 0$ and $a + b$ is divisible by 4
<b>6</b>	$\delta \geq 0$ and $b - a$ is divisible by 4
<b>7</b>	$\delta < 0$ and $a + b$ is divisible by 4
<b>8</b>	$\delta < 0$ and $b - a$ is divisible by 4

# Other optimizations

## CSA-related

- Repeatedly shifting
- Preserving sign
- Reducing number of CSAs
- Slower clock for carry-propagates

## Control overhead


- Precomputing control signals with high fanout
- Late select

## Performance Impact:

- **Number of carry-propagate adds is  $O(1)$  instead of  $O(\text{number of iterations})$**
- **Eliminate buffer delay to drive high fanout signals from critical path**



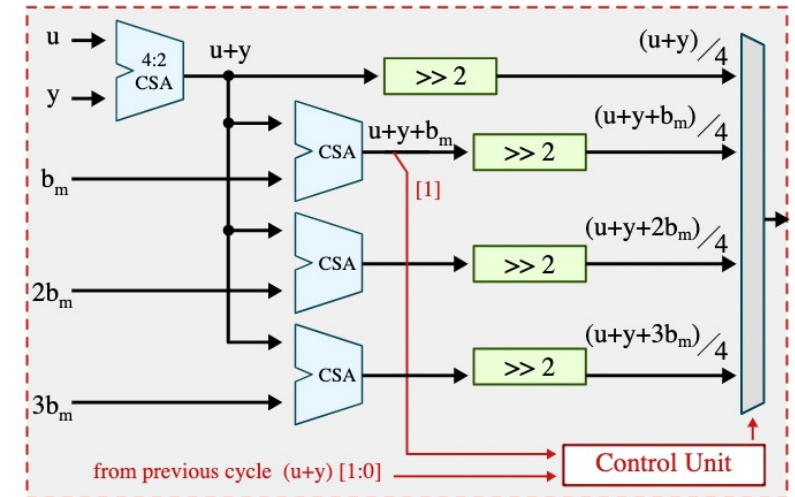
# SKY130 prototype with Efabless



Technology	SKY 130nm
Library	High density
Voltage	1.8 V
Frequency	325.4 MHz
Synthesized Area	5.66 mm <sup>2</sup>
Chip Area	10.28 mm <sup>2</sup>
Density	57%

# Critical Path: 3ns = 35 FO4 Inv Delay

Operation	Standard Cell	Delay (ns)	FO4 Inv Delay
DFF clk to Q	dfrtp_2	0.54	6.28
4:2 CSAs for $u + y$	fa_2	0.51	5.93
	fa_2	0.67	7.79
CSA for $(u + y) + 3 * b_m$	fa_2	0.45	5.23
	clkbuf_1	0.11	1.28
Half adder to correct +1	and2_4	0.13	1.51
	nor2_1	0.09	1.05
4:1 mux	mux2i_1	0.07	0.81
	a211oi_1	0.23	2.67
8:1 mux	nand2_1	0.07	0.81
	nand3_1	0.07	0.81
Library setup time	dfrtp_1	0.07	0.81
Total	N/A	3	35



# Area Breakdown

Module	Area ( $mm^2$ )	% of Area
Initial computation	0.27	4.8%
$a, b$ update (2-count)	0.31	5.5%
Bézout coefficient update (4-count)	3.84	68%
Control: $\delta, \alpha, \beta$ update	0.57	10%
Final result calculation	0.36	6.4%
JTAG	0.22	3.6%
Miscellaneous	0.10	1.7%
Total	5.66	100%



# Speedup achieved over prior work

Implementation	Technology	Algorithm	VDF Squaring Speedup*	Extended GCD Speedup
Zhu et al.'s VDF accelerator	28 nm	Chia's first competition round reference	10X	14X
Chia Network's C++ based on competition first round results	5 nm (M1)	NUDUPL	9X	67X

\*assuming only GCD is accelerated in hardware

# We can design efficient GCD accelerators for VDFs

- Verifiable delay functions are an exciting new cryptographic primitive with applications in blockchain
- Extended GCD requires 91% of the total VDF runtime
- Our design speeds up Chia Network's C++ by 9X

**VDFs can be sped up by an order of magnitude with our GCD hardware design techniques, laying the groundwork for fast hardware-accelerated VDFs in future blockchain systems.**