

# Introducing hardware generator debugger

and how we can debug at source-level

Keyi Zhang  
Computer Science Department  
Stanford University



# What are hardware generators?

First generation:

- “Preprocessor”-based string templating
- Using Perl or Python to preprocess the SystemVerilog to overcome the language limit.

# What are hardware generators?

First generation:

- “Preprocessor”
- Using C-like language

```
module `mname` (  
  clk, clk_en, reset,  
  //; for(my $i=0; $i<$pe_output_count; $i++) {  
    pe_output_`$i`,  
    //; }  
  //; for(my $i=0; $i<$sides; $i++) {  
    //; for(my $j=0; $j<$number_of_outputs; $j++) {  
      out_`$i`_`$j`,  
      in_`$i`_`$j`,  
      //; }  
    //; }  
  //; }
```

he

# What are hardware generators?

First generation:

- “Preprocessor”-based string templating
- Using Perl or Python to preprocess the SystemVerilog to overcome the language limit.

Second generation”

- Embedded domain specific languages that use host languages features to improve design productivity
- Chisel (scala), Manga (Python), PyMTL (Python)
- “Mini” compiler tool chains with IR

# What are hardware generators?

## First generation:

- “Preprocessor”
- Using a high-level hardware description language

```
class DoubleBufferFifo[T <: Data](gen: T, depth: Int) extends Fifo(gen: T, depth: Int) {
```

```
  private class DoubleBuffer[T <: Data](gen: T) extends Module {  
    val io = IO(new FifoIO(gen))
```

## Second generation:

- Embedded in a compiler
- Improved performance
- Chisel
- “Mini” compiler tool chains with IR

```
    val empty :: one :: two :: Nil = Enum(3)  
    val stateReg = RegInit(empty)  
    val dataReg = Reg(gen)  
    val shadowReg = Reg(gen)
```

```
    ...
```

# Google tried Chisel 3 years ago

**Code reuse.**

Software design constructs provide low-overhead approaches to refactoring common logic.

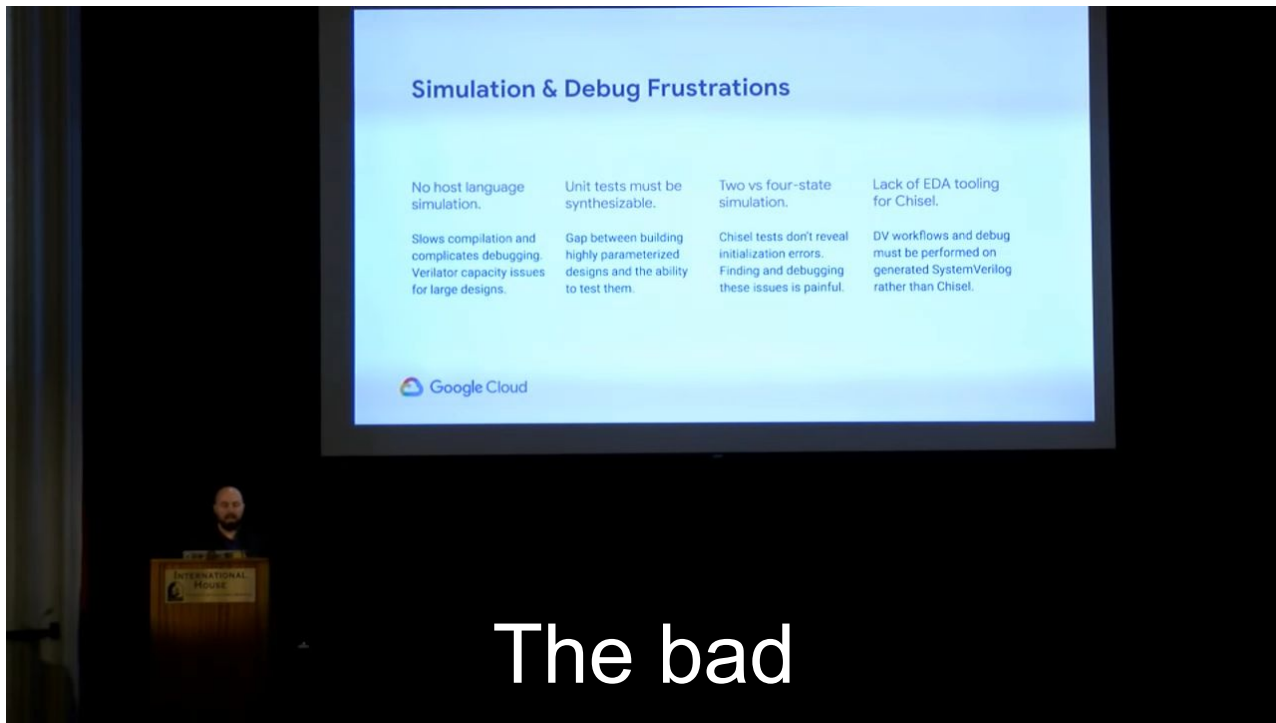
Functional refactoring allowed five distinct blocks in EdgeTPU to share 60% of their logic.

Google Cloud

```
graph TD; Common[Common] --> Unit1[Unit 1]; Common --> Unit2[Unit 2]; Common --> Unit3[Unit 3]; Common --> Unit4[Unit 4]; Common --> Unit5[Unit 5];
```

The good

# Google tried Chisel 3 years ago



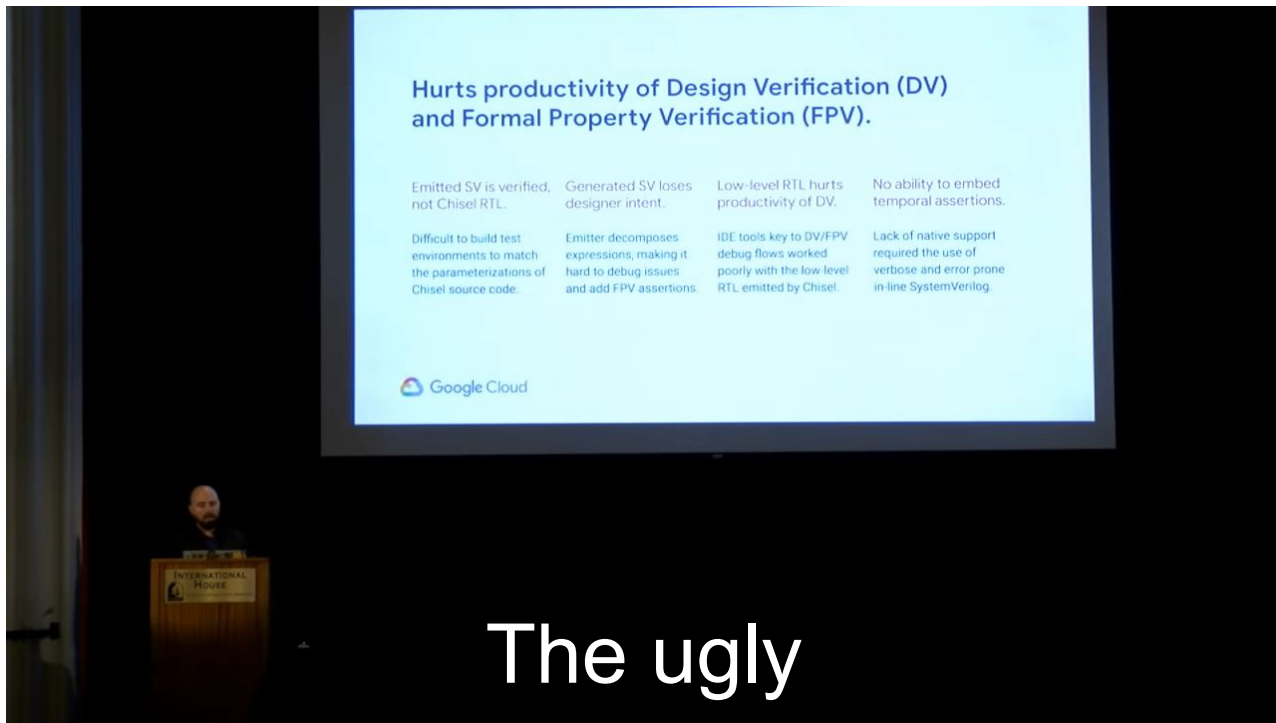
## Simulation & Debug Frustrations

No host language simulation.	Unit tests must be synthesizable.	Two vs four-state simulation.	Lack of EDA tooling for Chisel.
Slows compilation and complicates debugging. Verilator capacity issues for large designs.	Gap between building highly parameterized designs and the ability to test them.	Chisel tests don't reveal initialization errors. Finding and debugging these issues is painful.	DV workflows and debug must be performed on generated SystemVerilog rather than Chisel.

Google Cloud

# The bad

# Google tried Chisel 3 years ago





# Let's address the bad part: simulation and debugging

## Latest RocketChip + Chisel3

```
// SLT, SLTU
val slt =
  Mux(io.in1(xLen-1) === io.in2(xLen-1), io.adder_out(xLen-1),
    Mux(cmpUnsigned(io.fn), io.in2(xLen-1), io.in1(xLen-1)))
io.cmp_out := cmpInverted(io.fn) ^ Mux(cmpEq(io.fn), in1_xor_in2 === UInt(0), slt)

// SLL, SRL, SRA
val (shamt, shin_r) =
  if (xLen == 32) (io.in2(4,0), io.in1)
  else {
    require(xLen == 64)
    val shin_hi_32 = Fill(32, isSub(io.fn) && io.in1(31))
    val shin_hi = Mux(io.dw === DW_64, io.in1(63,32), shin_hi_32)
    val shamt = Cat(io.in2(5) & (io.dw === DW_64), io.in2(4,0))
    (shamt, Cat(shin_hi, io.in1(31,0)))
  }
val shin = Mux(io.fn === FN_SR || io.fn === FN_SRA, shin_r, Reverse(shin_r))
```

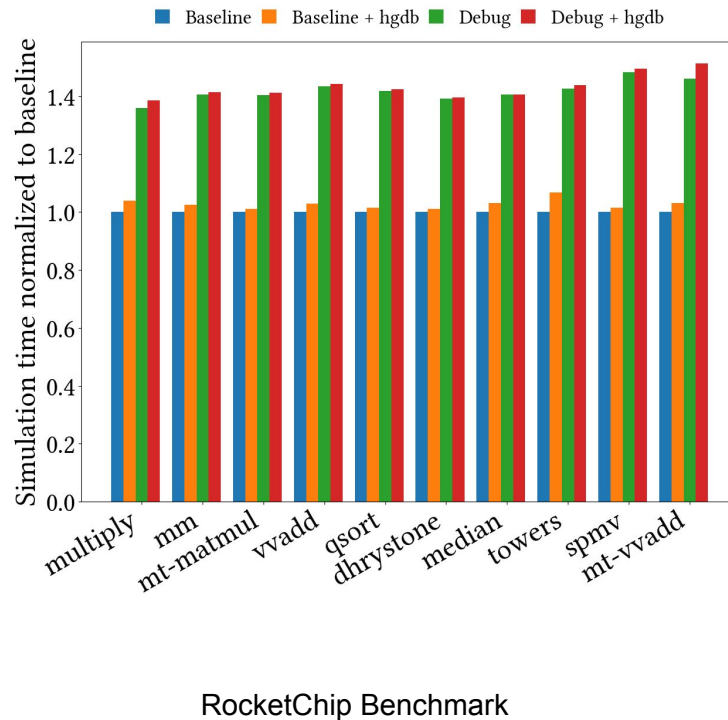
# Let's address the bad part: simulation and debugging

## Latest RocketChip + Chisel3


```
wire [63:0] _in2_inv_T_1 = ~io_in2; // @[ALU.scala 61:35]
wire [63:0] in2_inv = io_fn[3] ? _in2_inv_T_1 : io_in2; // @[ALU.scala 61:20]
wire [63:0] in1_xor_in2 = io_in1 ^ in2_inv; // @[ALU.scala 62:28]
wire [63:0] _io_adder_out_T_1 = io_in1 + in2_inv; // @[ALU.scala 63:26]
wire [63:0] _GEN_1 = {{63'd0}, io_fn[3]}; // @[ALU.scala 63:36]
wire _slt_T_7 = io_fn[1] ? io_in2[63] : io_in1[63]; // @[ALU.scala 68:8]
wire slt = io_in1[63] == io_in2[63] ? io_adder_out[63] : _slt_T_7; // @[ALU.scala 67:8]
wire _io_cmp_out_T_2 = ~io_fn[3]; // @[ALU.scala 44:26]
wire _io_cmp_out_T_4 = _io_cmp_out_T_2 ? in1_xor_in2 == 64'h0 : slt; // @[ALU.scala 69:41]
wire _T_2 = io_fn[3] & io_in1[31]; // @[ALU.scala 76:46]
wire [31:0] _T_4 = _T_2 ? 32'hffffffff : 32'h0; // @[Bitwise.scala 72:12]
wire [31:0] hi = io_dw ? io_in1[63:32] : _T_4; // @[ALU.scala 77:24]
wire hi_1 = io_in2[5] & io_dw; // @[ALU.scala 78:33]
wire [4:0] lo = io_in2[4:0]; // @[ALU.scala 78:60]
wire [5:0] shamt = {hi_1, lo}; // @[Cat.scala 30:58]
wire [31:0] lo_1 = io_in1[31:0]; // @[ALU.scala 79:34]
wire [63:0] shin_r = {hi, lo_1}; // @[Cat.scala 30:58]
wire _shin_T_2 = io_fn == 4'h5 | io_fn == 4'hb; // @[ALU.scala 81:35]
wire [63:0] _shin_T_6 = {{32'd0}, shin_r[63:32]}; // @[Bitwise.scala 103:31]
wire [63:0] _shin_T_8 = {shin_r[31:0], 32'h0}; // @[Bitwise.scala 103:65]
wire [63:0] _shin_T_10 = _shin_T_8 & 64'hffffffff00000000; // @[Bitwise.scala 103:75]
wire [63:0] _shin_T_11 = _shin_T_6 | _shin_T_10; // @[Bitwise.scala 103:39]
wire [63:0] _GEN_2 = {{16'd0}, _shin_T_11[63:16]}; // @[Bitwise.scala 103:31]
wire [63:0] _shin_T_16 = _GEN_2 & 64'hffff0000ffff; // @[Bitwise.scala 103:31]
wire [63:0] _shin_T_18 = {_shin_T_11[47:0], 16'h0}; // @[Bitwise.scala 103:65]
wire [63:0] _shin_T_20 = _shin_T_18 & 64'hffff0000ffff0000; // @[Bitwise.scala 103:75]
wire [63:0] _shin_T_21 = _shin_T_16 | _shin_T_20; // @[Bitwise.scala 103:39]
wire [63:0] _GEN_3 = {{8'd0}, _shin_T_21[63:8]}; // @[Bitwise.scala 103:31]
```

# hgdb: batteries included

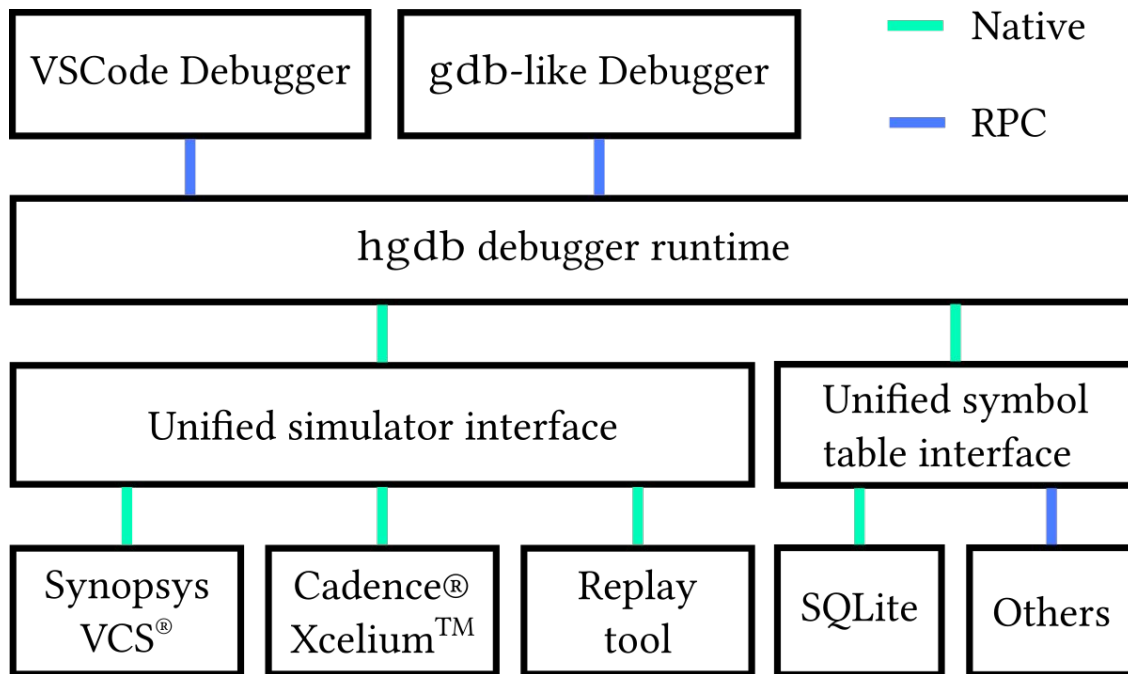
- Source-level debugging
- Minimal performance overhead
- No RTL changes required <sup>1</sup>
- Two complete debuggers
  - VSCode
  - gdb-inspired console debugger
- All major simulators
  - Big 3
  - Verilator
  - iverilog
- FSDB and VCD Replay
  - Reverse debugging!



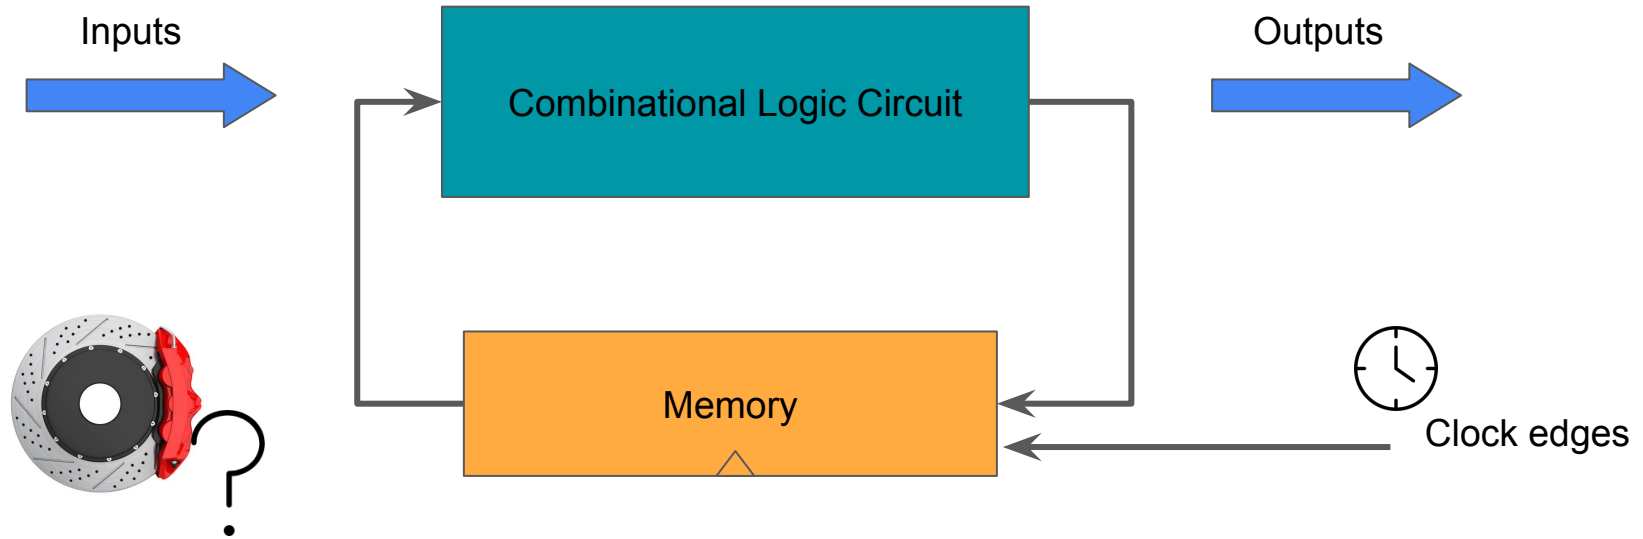
# What hgdb brings to the table



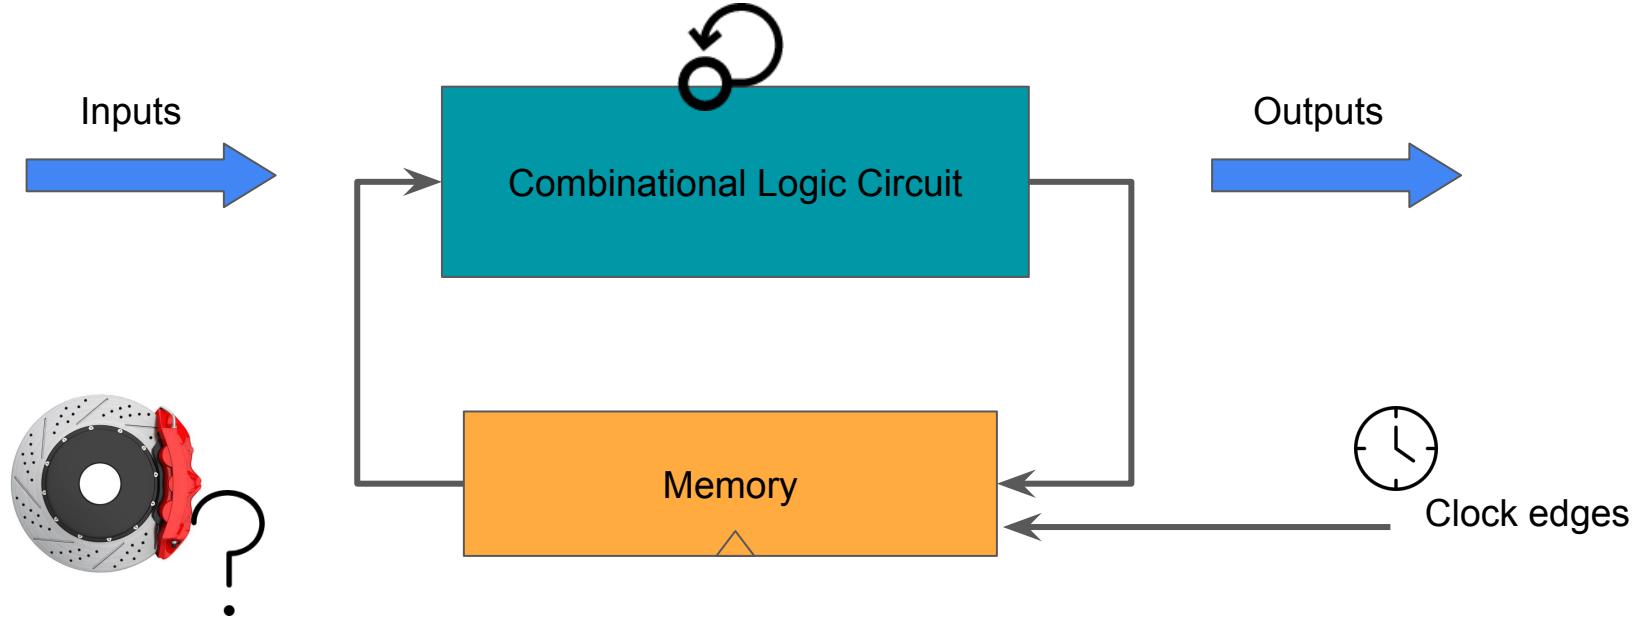
# How hgdb works



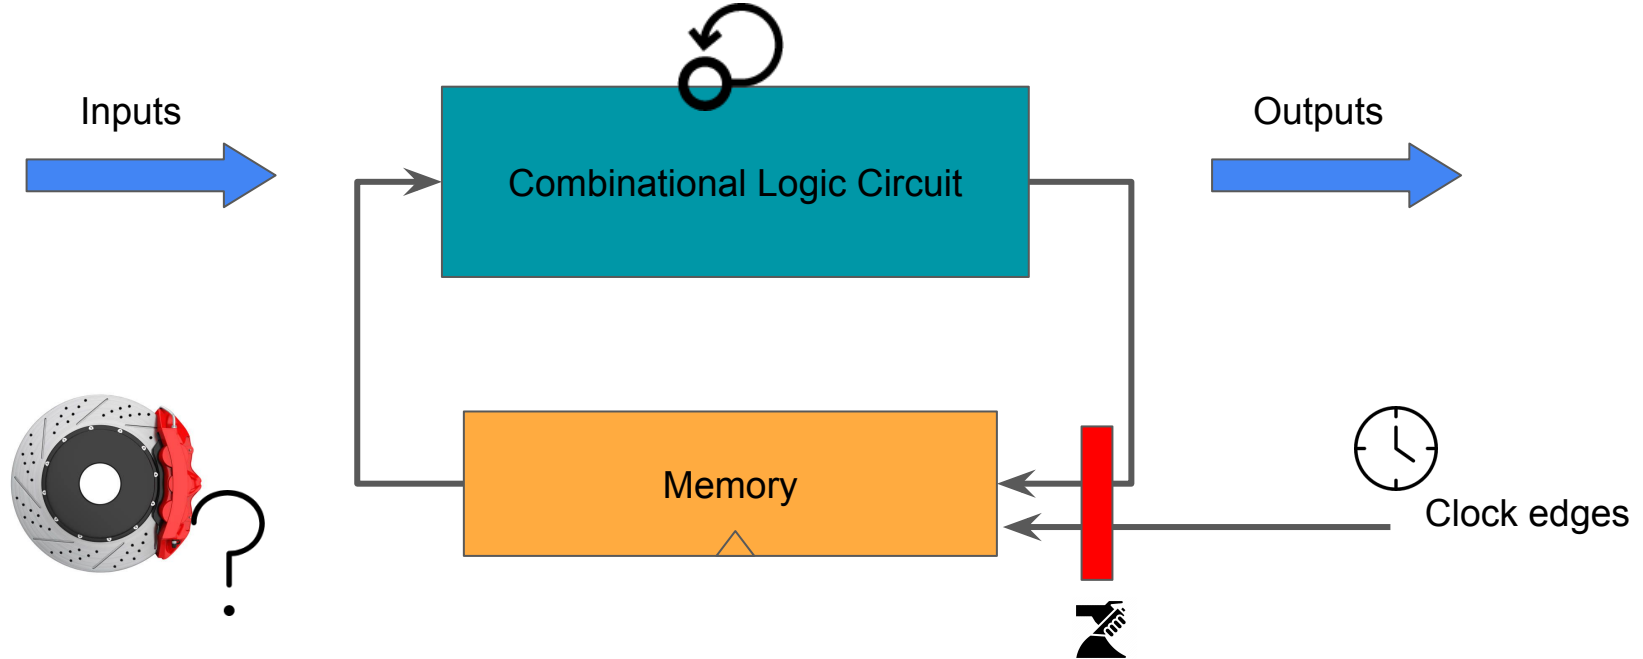
# Keeps the overhead low: breakpoint emulation



# Keeps the overhead low: breakpoint emulation



# Keeps the overhead low: breakpoint emulation






# But does it work all the time?

No. Unfortunately. 😞

```
logic [31:0] sum;  
logic [31:0] data[4];
```

```
always_comb begin  
    sum = 0;  
    for (int i = 0; i < 4; i++) begin  
        if (data[i] % 2) begin  
            sum += data[i];  
        end  
    end  
end
```

Combinational logic



Partial results lost when the state stabilizes



# Loop unrolling and SSA to rescue

↖ Static single assignment

```
logic [31:0] sum;
logic [31:0] data[4];

always_comb begin
    sum = 0;
    if (data[0] % 2) sum += data[0];
    if (data[1] % 2) sum += data[1];
    if (data[2] % 2) sum += data[2];
    if (data[3] % 2) sum += data[3];
end
```

Loop unroll

```
logic [31:0] sum, sum0, sum1, sum2, sum3;
logic [31:0] data[4];

assign sum0 = data[0] % 2 ? data[0] : 0;
assign sum1 = data[1] % 2 ? sum0 + data[1] : sum0;
assign sum2 = data[2] % 2 ? sum1 + data[2] : sum1;
assign sum3 = data[3] % 2 ? sum2 + data[3] : sum2;
assign sum = sum3;
```

SSA Transform

# Loop unrolling and SSA to rescue

```
logic [31:0] sum;  
logic [31:0] data[4];
```

```
always_comb begin  
    sum = 0;  
    for (int i = 0; i < 4; i++) begin  
        if (data[i] % 2) begin  
            sum += data[i];  
        end  
    end  
end
```

Original form

```
logic [31:0] sum, sum0, sum1, sum2, sum3;  
logic [31:0] data[4];
```

```
assign sum0 = data[0] % 2 ? data[0] : 0;  
assign sum1 = data[1] % 2 ? sum0 + data[1] : sum0;  
assign sum2 = data[2] % 2 ? sum1 + data[2] : sum1;  
assign sum3 = data[3] % 2 ? sum2 + data[3] : sum2;  
assign sum = sum3;
```

Final form

Multiple mapping!

sum



# Loop unrolling and SSA to rescue

```
logic [31:0] sum;  
logic [31:0] data[4];
```

```
always_comb begin  
    sum = 0;  
    for (int i = 0; i < 4; i++) begin  
        if (data[i] % 2) begin  
            sum += data[i];  
        end  
    end  
end
```

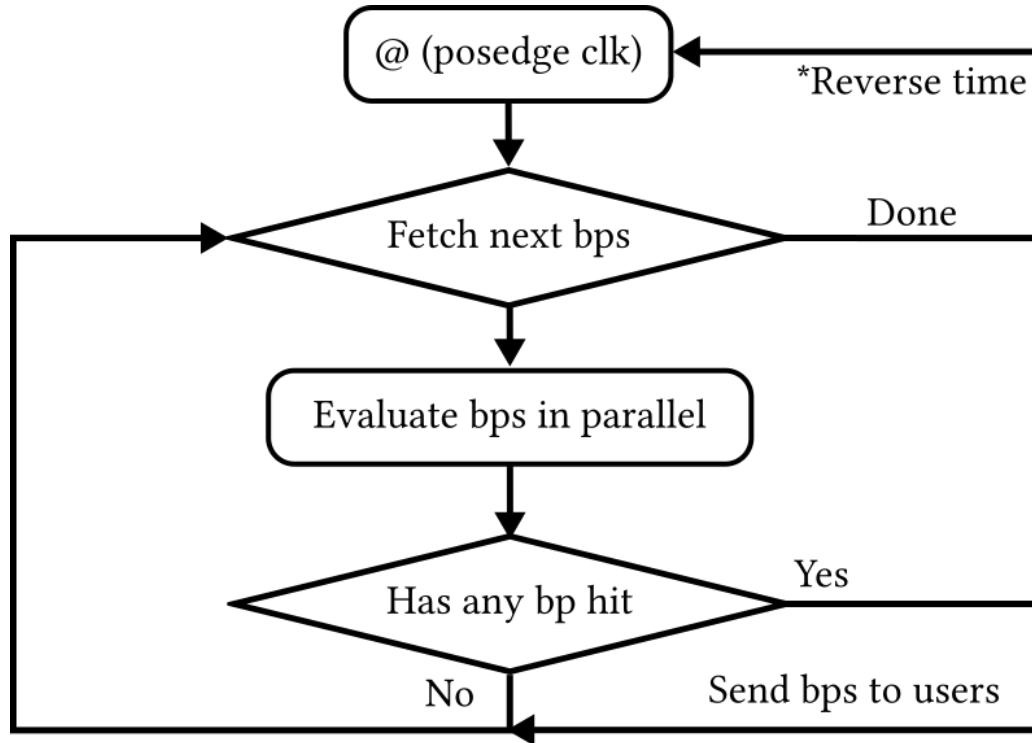
Original form

```
logic [31:0] sum, sum0, sum1, sum2, sum3;  
logic [31:0] data[4];
```

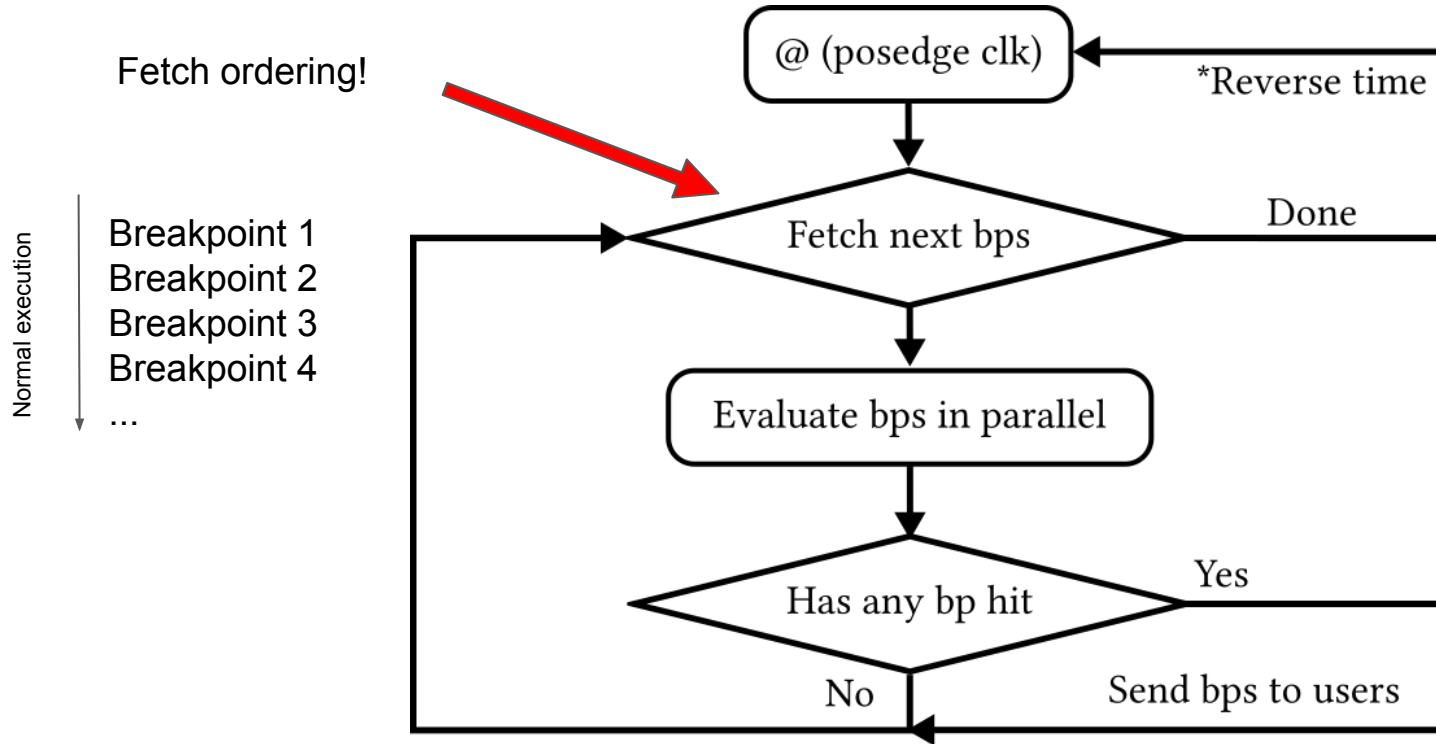
```
assign sum0 = data[0] % 2 ? data[0] : 0;  
assign sum1 = data[1] % 2 ? sum0 + data[1] : sum0;  
assign sum2 = data[2] % 2 ? sum1 + data[2] : sum1;  
assign sum3 = data[3] % 2 ? sum2 + data[3] : sum2;  
assign sum = sum3;
```

Final form

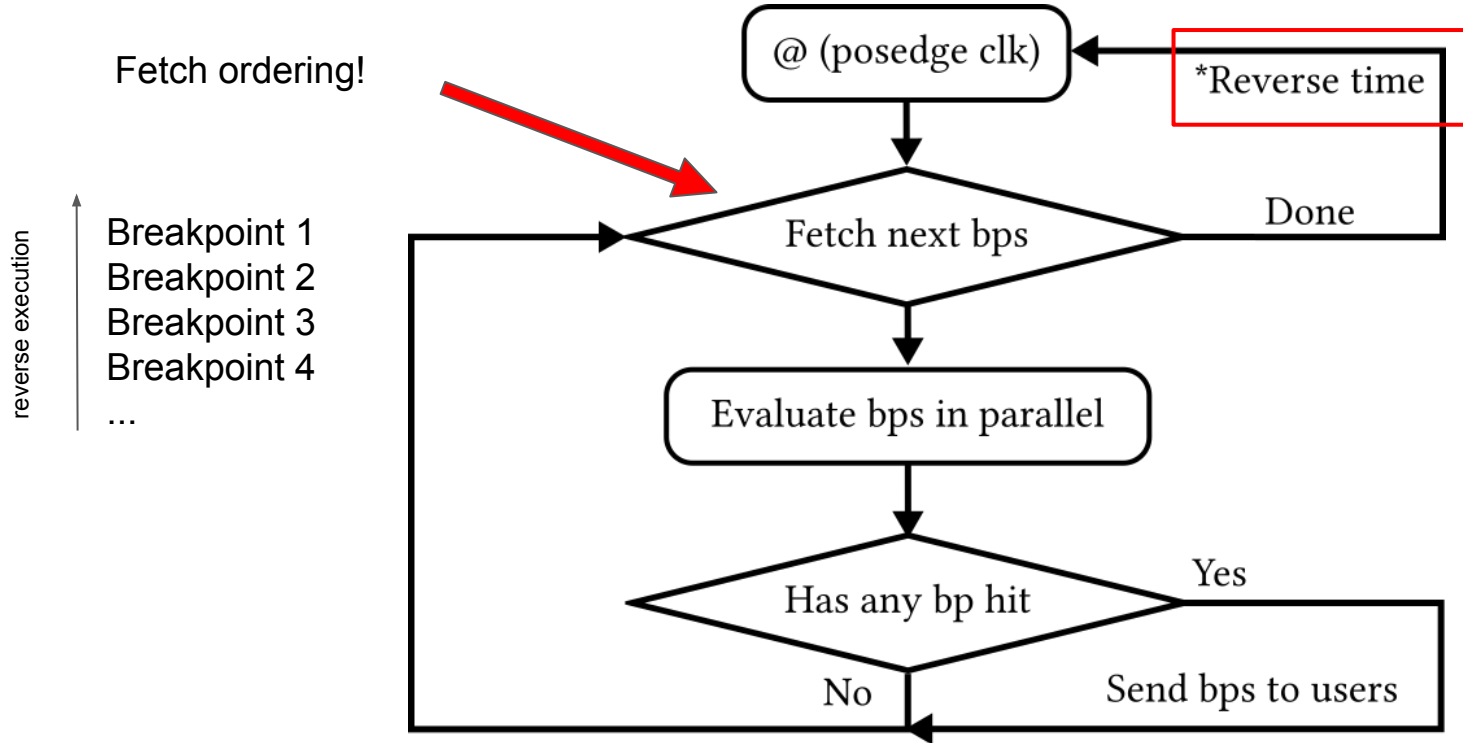
# Breakpoint emulation loop



# Breakpoint emulation loop



# Breakpoint emulation loop



# Unified simulator interface

## Primitives:

- Place callback on value change
- Get signal values
- Get design hierarchy
- Reverse time (optional)

db-like Debugger

Native

RPC

to debugger runtime

Unified simulator interface

Unified symbol  
table interface

Synopsys  
VCS®

Cadence®  
Xcelium™

Replay  
tool

SQLite

Others



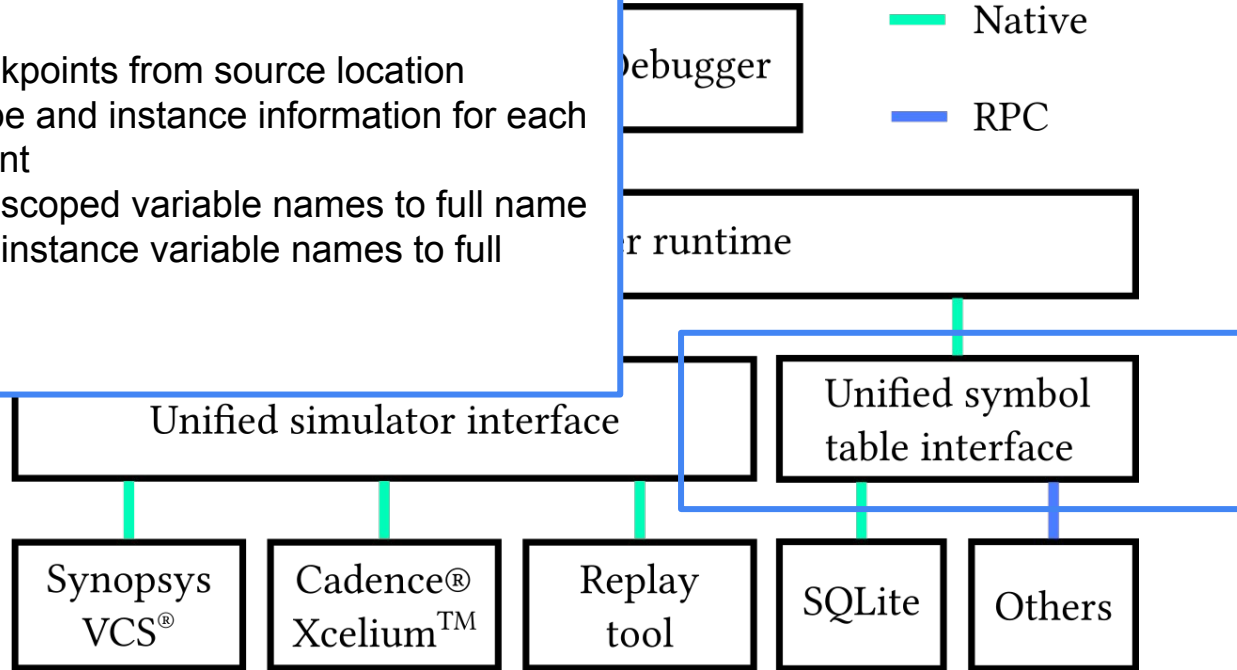
enjenir work!



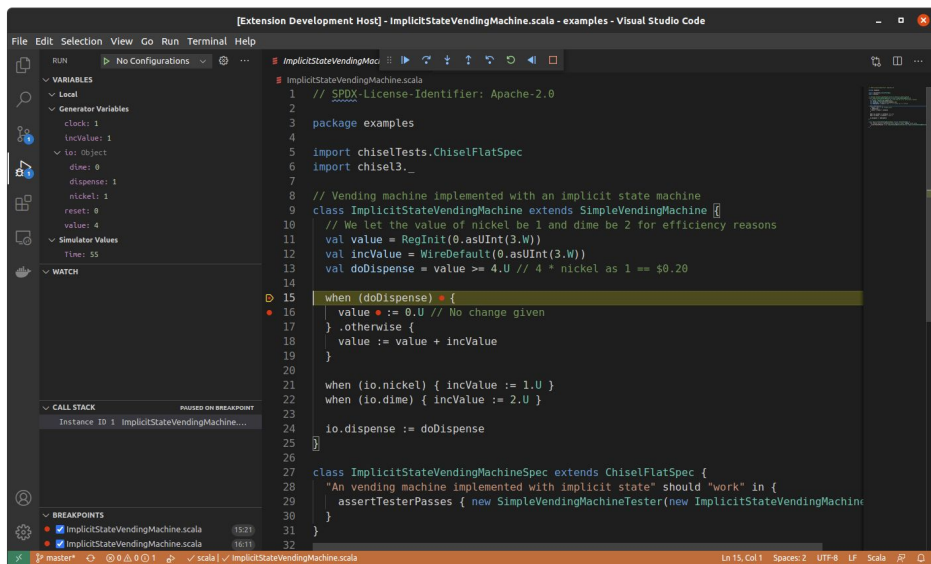
# Unified symbol table interface

## Primitives:

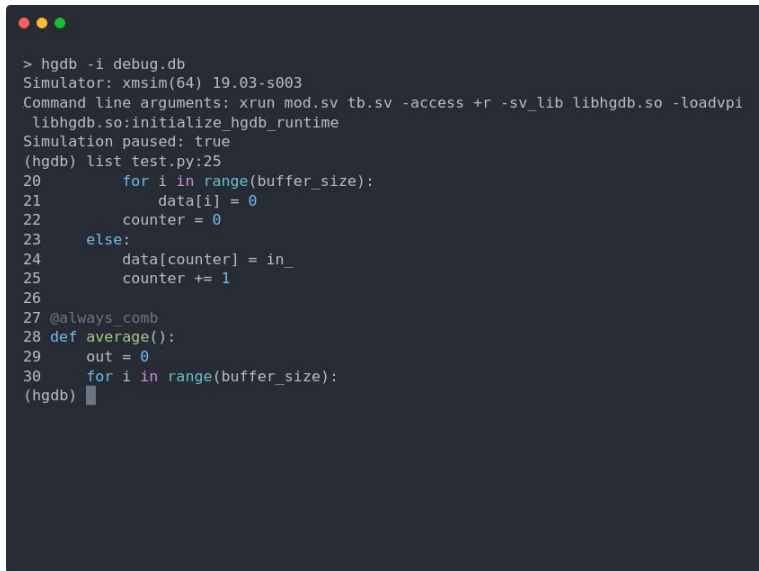
- Get breakpoints from source location
- Get scope and instance information for each breakpoint
- Resolve scoped variable names to full name
- Resolve instance variable names to full name



# Hgdb debuggers

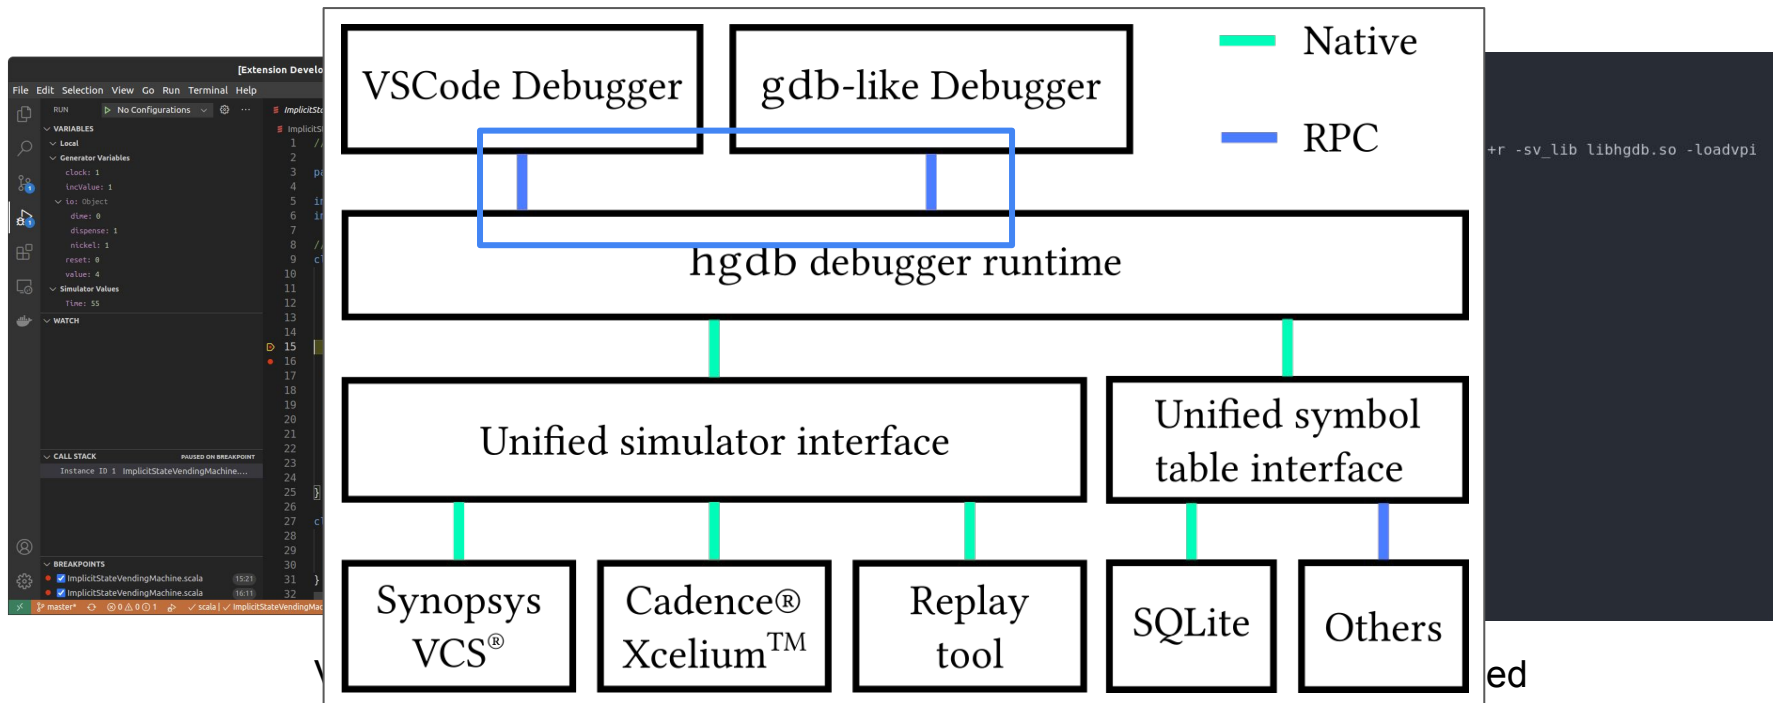


Visual Studio Code



Terminal based

# Hgdb debuggers



# How hgdb works with Chisel



```
node _T = bits(io.fn, 3, 3) @[ALU.scala 40:29]
node _T_1 = bits(io.in1, 31, 31) @[ALU.scala 76:55]
node _T_2 = and(_T, _T_1) @[ALU.scala 76:46]
node _T_3 = bits(_T_2, 0, 0) @[Bitwise.scala 72:15]
node _T_4 = mux(_T_3,
                UInt<32>("h0ffffffff"),
                UInt<32>("h00")) @[Bitwise.scala 72:12]
node _T_5 = eq(io.dw, UInt<1>("h01")) @[ALU.scala 77:31]
node _T_6 = bits(io.in1, 63, 32) @[ALU.scala 77:48]
node hi = mux(_T_5, _T_6, _T_4) @[ALU.scala 77:24]
```

# How hgdb works with Chisel



```
node _io_adder_out_T = add(io.in1, in2_inv) @[ALU.scala 63:26]
node _io_adder_out_T_1 = tail(_io_adder_out_T, 1) @[ALU.scala 63:26]
node _io_adder_out_T_2 = bits(io.fn, 3, 3) @[ALU.scala 40:29]
node _io_adder_out_T_3 = add(_io_adder_out_T_1, _io_adder_out_T_2)
@[ALU.scala 63:36]
node _io_adder_out_T_4 = tail(_io_adder_out_T_3, 1) @[ALU.scala 63:36]
io.adder_out <= _io_adder_out_T_4 @[ALU.scala 63:16]
```



```
wire [63:0] _io_adder_out_T_1 = io.in1 + in2_inv; // @[ALU.scala 63:26]
wire [63:0] _GEN_1 = {{63'd0}, io.fn[3]}; // @[ALU.scala 63:36]
assign io_out = ~io_dw ? _io_out_T_2 : out; // @[ALU.scala 96:28 ALU.scala 96:37 ALU.scala 93:10]
assign io_adder_out = _io_adder_out_T_1 + _GEN_1; // @[ALU.scala 63:36]
```

# How hgdb works with Chisel



AND THE  
UGLY



Let's insert DontTouchAnnotation everywhere!

```
mkdir -p /home/keyi/workspace/rocket-chip-debug/vsim/generated-src/  
/home/keyi/workspace/hgdb-firrtl/bin/firrtl -i /home/keyi/workspace/rocket-chip-debug/vsim/generat  
ed-src/freechips.rocketchip.system.DefaultConfig.fir \  
-o /home/keyi/workspace/rocket-chip-debug/vsim/generated-src/freechips.rocketchip.system.Defau  
ltConfig.v \  
-X verilog \  
--infer-rw TestHarness \  
--hgdb-toml rocket-chip.toml \  
-O 0 \  
--repl-seq-mem -c:TestHarness:-o:/home/keyi/workspace/rocket-chip-debug/vsim/generated-src/fre  
echips.rocketchip.system.DefaultConfig.conf \  
-faf /home/keyi/workspace/rocket-chip-debug/vsim/generated-src/freechips.rocketchip.system.Def  
aultConfig.anno.json \  
-td /home/keyi/workspace/rocket-chip-debug/vsim/generated-src/freechips.rocketchip.system.Defau  
ltConfig/ \  
-fct firrtl.passes.InlineInstances, \
```

```
Exception in thread "main" firrtl.transforms.DontTouchAnnotation$DontTouchNotFoundException: Targe  
t marked dontTouch (InterruptBusWrapper.auto_int bus_int_out_0) not found!  
It was probably accidentally deleted. Please check that your custom transforms are not responsible  
and then  
file an issue on GitHub: https://github.com/freechipsproject/firrtl/issues/new
```

# It finally works!

**Input:** *CircuitState*

**Output:** *Table*

*Annotations*  $\leftarrow \{\}$ ;

**foreach** *node*  $\in$  *CircuitState* **do** // First pass

**if** *node* is statement **then**

        | *node.enable*  $\leftarrow$  *ComputeEnableCondition*(*node*);

**end**

*Annotation*  $\leftarrow$  *Annotations*  $\cup$  {*node*}

**end**

// FIRRTL transformations;

*IRNodes*  $\leftarrow \{\}$ ;

**foreach** *node*  $\in$  *Annotations* **do**

**if** *node*  $\in$  *CircuitState* **then**

        | *IRNodes*  $\leftarrow$  *IRNodes*  $\cup$  *node*;

**end**

**end**

*Table*  $\leftarrow$  *ComputeSymbolTable*(*IRNodes*);

First pass:

- Insert annotation and compute enable condition
- (Debug mode) insert DontTouchAnnotation

Second pass:

- Collect annotations and only compute symbol table if the IRNode still exists



<https://github.com/Kuree/hgdb-firrtl>

# It finally works!

**Input:** *CircuitState*

**Output:** *Table*

*Annotations*  $\leftarrow \{\}$ ;

**foreach** *node*  $\in$  *CircuitState* **do** //

**if** *node* is statement **then**

*node.enable*  $\leftarrow$  *ComputeEnc*

**end**

*Annotation*  $\leftarrow$  *Annotations*  $\cup$  {*i*}

**end**

// FIRRTL transformations;

*IRNodes*  $\leftarrow \{\}$ ;

**foreach** *node*  $\in$  *Annotations* **do**


**if** *node*  $\in$  *CircuitState* **then**

*IRNodes*  $\leftarrow$  *IRNodes*  $\cup$  *node*

**end**

**end**


*Table*  $\leftarrow$  *ComputeSymbolTable*(*IRNodes*);



compute enable

ontTouchAnnotation

only compute  
ode still exists



<https://github.com/Kuree/hgdb-firrtl>



# Future work in progress

- Work presented to LLVM/CIRCT and SiFive
  - Integrate hgdb to CIRCT system
  - Leverage LLVM compiler toolchain
- Watchpoint
  - We believe watchpoint can be effectively emulated with extra information from the IR
  - Similar to breakpoint emulation
  - Can replace driver tracing if used with reverse debugging
- Symbol table redesign
  - Current implementation “too flat”.
  - More akin to DWARF but simpler and easy to read/write

# Conclusion

- Hgdb bridges the gap between hardware generator frameworks and existing simulators
  - Efficient
  - Works with all major simulator vendors
  - What you debug is what you tape out
- Hardware generators are fairly new, and debugging infrastructure is almost non-existing.
  - CIRCIIT is the future???
- Hgdb is an open-source framework
  - Contributions are welcome!



<https://github.com/Kuree/hgdb>



Thank you